

Measuring Appellate Reason-Giving in Saudi Commercial Courts

Abdullah Almohammedi

Independent Researcher

abdullah.m.almohammedi@gmail.com

ORCID: 0009-0001-0832-0995

Abstract

Saudi commercial judgments are published with the appellate decision attached to the judgment it reviews, giving 13,924 matched pairs of the same dispute at two levels. 74.4 per cent of appellate judgments give no reasons of their own, affirming on the reasons below; whether a circuit writes is not independent of what it decides, rising from 26.2 per cent when it affirms to 34.3 when it reverses. Of appeals reaching the merits, 10.8 per cent disturb the judgment. A logistic model on the 9,563 decided between 1439 and 1444 — the window in which the record is mature enough to read — finds the procedural share of a judgment's own citations the strongest protection against disturbance (odds ratio 0.74) and one of three terms of eleven to survive correction for multiple testing. The judgment's citation count and whether it wrote its own reasons predict nothing; its length predicts weakly, and in the direction opposite to the one a reader would guess.

Keywords: appellate review; judicial reason-giving; empirical legal studies; court administration; Saudi Arabia

1 Introduction

An appeal is studied where appellate decisions are digitised, which in practice has meant the United States and a handful of other common-law systems. The second instance elsewhere is known largely through court statistics: how many appeals were brought, how many succeeded. What the appellate court actually wrote, set beside what the court below wrote in the same dispute, is rarely available at scale in any jurisdiction, because the two documents live in different places.

In Saudi Arabia they live in the same record. The Ministry of Justice publishes commercial judgments in full text, and where a case went up, the appellate decision is published attached to the judgment it reviews. That yields 13,924 matched pairs — the same dispute seen twice — inside a corpus of 50,666 judgments, with a further 1,459 appellate decisions published alone. Nothing about the pairing was designed for research; it is a by-product of how the publisher stores a case file. But it removes the objection that defeats most comparisons between levels, which is that the cases differ. Here they do not differ. They are the same cases.

Three things follow, and the first is about the institution rather than any case. 74.4 per cent of appellate judgments in this corpus give no reasons of their own: they affirm on the reasons below, in the formula *mahmulan 'ala asbabih*. The second instance therefore adds its own statement of reasoning to the public record in about one appeal in four. Whether it does so is not independent of what it decides — circuits

wrote their own reasons in 26.2 per cent of affirmances and 34.3 per cent of reversals — which matters both as a finding and as a constraint on the third result below.

Second, of the appeals that reached the merits, 10.8 per cent disturbed the judgment: reversed it, replaced it with a fresh judgment, or varied it. The rest affirmed. A further 14.7 per cent of all appellate decisions never reached the merits at all, failing at the threshold, and are excluded from that rate rather than counted as affirmances, because an inadmissible objection says nothing about whether the judgment below was sound.

Third, a logistic model asks what separates the judgments that were disturbed from those that survived, using only what can be read from the first-instance document. On the 9,563 decided between 1439 and 1444, the strongest protection is the procedural share of a judgment's own statutory citations (odds ratio 0.74), one of only three terms of eleven to survive correction for multiple testing. A judgment resting on the law of procedure is harder to disturb than one resting on the law of obligations. The number of instruments it cites and whether it wrote its own reasons predict nothing; its length predicts a little, and longer judgments are disturbed slightly more often rather than less.

Three methodological points are reported rather than assumed. The appellate disposition is read from the operative part of the appellate document and nowhere else, because an appellate recital narrates the history of the case and a pattern applied to the whole document reads an earlier affirmance as this court's ruling; the classifier was corrected four times against hand-read samples before it was used. And the year in which a judgment was given cannot be read across the whole span: a recent judgment appears here with an appellate decision only if the appeal was both brought and decided before the corpus was collected, so the disturbed rate falls from 11.8 per cent inside the mature window to almost nothing in the last years of the span. The window is therefore an argument to the model rather than a silent assumption, and the article reports the mature one.

The third is about the unit of analysis. Where a comparison is between two levels of the same dispute, it is made over pairs and not over citations: citations cluster inside judgments, and pooling them weights a dispute by how many citations it happens to carry. That distinction is not cosmetic here. It reverses the sign of one of this article's own earlier results, which §5 reports and withdraws.

2 Related work

What an appeal is for. The economic account treats the appeal as a device for correcting error at a cost, in which the party who lost has private information about whether the decision was wrong and the system relies on that information to select the cases worth reviewing (Shavell 1995). The account predicts a low disturbance rate among the appeals that are brought, since litigants who know they were rightly beaten should not appeal; what it does not predict is which surviving judgments differ from those overturned, which is an empirical question this corpus can answer directly.

Measuring appellate outcomes. The empirical literature on appeal rates and outcomes is built largely on United States federal and state court statistics, and is centrally concerned with who wins on appeal and why plaintiffs do worse there than at trial (Eisenberg 2004). That work reads outcomes from docket records: the appellate decision is a category rather than a text. This article reads the text, which permits questions a docket cannot answer — whether the appellate court reasoned at all, and what it cited when it

did — at the price of being confined to what one publisher publishes.

Reason-giving, and its limits. That courts should give reasons is among the least controversial propositions in procedural theory, and among the least absolute in practice. Cohen (2015) sets out the values that pull against it — efficiency, sincerity, and the risk of guidance a court is not ready to give — and shows common-law and civil-law systems converging on similar compromises. Affirming on the reasons below is one such compromise, and a defensible one: reasons adopted are reasons given. Whether a court adopts or writes is nevertheless observable, and the rate at which the second instance writes has not been measured for this jurisdiction. It is measured here.

The filters between a dispute and this dataset. Two selection problems bound everything below and cannot be removed by better extraction. Cases that settle never produce a judgment, and settled disputes differ systematically from tried ones (Priest and Klein 1984); of the judgments given, publication is a second filter, and published and unpublished decisions have been shown to differ in kind rather than merely in number (Siegelman and Donohue 1990). A third filter is specific to appeals and is the more binding here: a judgment enters this analysis only if it was appealed. Nothing in what follows describes judgments that no one challenged.

The jurisdiction. Scholarship on Saudi law has been overwhelmingly doctrinal and institutional, and until recently had little else to work with; Vogel (2000) remains the standard English-language treatment of the courts, written before judgments were published at scale. The commercial courts whose work this article measures were reorganised under a new law in 2020, and the corpus is dominated by the years since.

3 Data

3.1 A record that carries both levels

The Ministry of Justice publishes Saudi commercial judgments in full text through the interface its own portal uses. A published record holds the judgment, and where the case was appealed it holds the appellate decision as well, in a separate field of the same record. 15,383 records carry an appellate decision; 13,924 of those carry it beside the first-instance judgment, and 1,459 carry the appellate decision alone.

That storage decision is the whole of this article’s data advantage, and it was not made for research. It also caused a defect worth stating, because it shows what the pairing does to naive processing: an earlier analysis of this corpus segmented each record on its own headings, and where a record held two judgments the segmentation ran from the reasons of the first into the disposition of the second. Each document is now segmented separately.

3.2 What the published record is

The corpus holds 50,666 judgments, overwhelmingly commercial and concentrated in three cities and the last few Hijri years. Section 6 states the consequences. Two features of the published fields matter for method rather than for interpretation. The interface exposes `judgmentFacts`, `judgmentReasons` and `judgmentRuling`, which are null in every record, so the structure of a judgment must be read

from its text. And it exposes `isAppeal`, which is present, well-typed, and false on every record — including every record that carries an appellate decision. First instance and appeal are separated here by the fields that hold the text, not by the flag provided for the purpose.

4 Method

4.1 Reading the appellate disposition

The disposition is read from the operative part of the appellate document and from nowhere else. An appellate recital routinely narrates the history of the case, including earlier affirmances, so a pattern applied to the whole document reports an old affirmance as this court’s ruling. The operative part is located by the markers Saudi circuits use to introduce it, and only the text after the last such marker is classified; where the operative part quotes the judgment below — which it usually does, at length — the quoted disposition is cut away before classification, and the label is whichever disposition the court states first.

Dispositions are classified as affirmed, reversed, replaced by a fresh judgment, varied, not admitted, disposed of on another ground, or unclear. Nothing is forced: a document whose operative part cannot be located, or whose disposition matches none of these, is unclear and is reported as such.

The classifier was corrected four times, each time against a hand-read sample rather than by inspection. The operative part is fully vowelised in a few hundred judgments and a plain-script pattern missed all of them; a refusal of a request for arbitration is not an inadmissible appeal; the window swallowed the quoted judgment below; and a fixed order of tests read *affirming a judgment that had reversed an administrative decision* as a reversal. On a random sample of twenty classifications read after the last correction, none was wrong.

4.2 Whether the appeal reasoned

Saudi judgments carry their own structure — *al-waqa’i*, then *al-asbab*, then the operative part — and an appellate document is counted as having reasoned when that structure is present within it. This is a count of whether the second instance put its own reasoning on the record, not a judgment about quality: a circuit affirming on the reasons below has given reasons by adoption.

4.3 Modelling disturbance

The model is a logistic regression of disturbance on features of the first-instance judgment. Disturbed means reversed, replaced or varied; affirmed means affirmed. Appeals that never reached the merits are excluded rather than coded as affirmances.

Every feature is read from the first-instance document alone: the city, the Hijri year, the length of the text, the number of statutory citations, the number of distinct instruments cited, the procedural share of those citations, whether the judgment carried its own reasons under the headings, whether its operative part disposed of the case on the merits rather than on a preliminary point, and whether a preliminary defence was raised in it. A model that saw the appellate text would be predicting the outcome from itself.

The fit is by iteratively reweighted least squares, written out rather than imported: the replication guide promises that the analysis runs with no third-party package, and a reviewer running the pipeline

Table 1: Appellate dispositions across 15,383 published appellate decisions.

Disposition	Judgments	Share
Affirmed	10,021	65.1%
Reversed	736	4.78%
Replaced by a fresh judgment	478	3.11%
Varied	3	—
Not admitted	2,261	14.7%
Disposed of on another ground	435	2.83%
Unclear	1,449	9.42%

should not discover otherwise because one script wanted a matrix inverse.

5 Results

5.1 What appeals do

Table 1 gives the dispositions. 11,238 decisions reached the merits, and 10.8 per cent of those disturbed the judgment below. The rate is low, as the error-correction account predicts of appeals that are brought at all, and the striking figure beside it is the 14.7 per cent that failed at the threshold: about one appellate decision in seven never reached the question the appellant wanted answered.

5.2 Reason-giving at the second instance

3,931 appellate judgments — 25.6 per cent — carry reasons of their own. The remaining 74.4 per cent affirm or dispose on the reasons below.

Whether a circuit writes is not independent of what it decides. It wrote in 26.2 per cent of affirmances, 34.3 per cent of reversals, 33.2 per cent where it replaced the judgment with its own, and 17.4 per cent where the appeal failed at the threshold. A circuit that leaves the judgment standing mostly adopts the reasoning below; a circuit that disturbs it mostly explains itself. That is what one would want of an appellate court, and it is also what makes the comparison in the next subsection a biased sample of appeals rather than a sample of appeals.

5.3 Reasoning compared, where both levels reasoned

Both levels reasoned in 3,106 pairs, 22.3 per cent of the pairs. Pooled over citations, the appellate reasons look more procedural than the reasons below — 95.6 per cent against 93.2 — and drawn from a narrower set of instruments, 29 distinct instruments against 42.

That figure is computed on the wrong unit, and it does not survive being put on the right one. Citations are not independent draws: they cluster inside judgments, so pooling them weights a pair by how many citations it happens to carry, and a few citation-heavy appellate documents can carry the average. Table 2 runs the same comparison over pairs. Each of the 876 pairs in which both levels cite a matched instrument contributes one number — the procedural share of the appellate reasons minus the procedural share of the reasons below — and the null that the levels do not differ is tested by flipping the sign of that number at random, which is the exact randomisation test for paired data and assumes nothing about its distribution.

Table 2: The same comparison over pairs rather than over citations. A negative difference means the appellate reasons are less procedural than the reasons below. Two-sided p by sign-flip permutation.

Pairs where both levels cited	Pairs	Mean difference	p
All	876	−1.5	.087
The appeal affirmed	417	−1.2	.280
The appeal disturbed	135	−7.1	.033

Table 3: Logistic regression of disturbance on features of the first-instance judgment, 1439–1444. 9,563 judgments, base rate 11.8 per cent.

Feature of the judgment below	Odds ratio	p	Holm p
Procedural share of its citations	0.74	<.001	.006
Hijri year	1.30	<.001	<.001
Dammam	1.44	.003	.029
Riyadh	1.09	.401	1.000
Jeddah	1.09	.438	1.000
Disposed on the merits	0.86	.020	.157
Preliminary defence raised	1.13	.156	.778
Wrote its own reasons	0.85	.117	.699
Length (log characters)	1.11	.044	.307
Number of citations	0.98	.509	1.000
Distinct instruments cited	1.05	.485	1.000

The sign reverses. Over all pairs the mean paired difference is −1.5 percentage points — the appellate reasons are, if anything, *less* procedural than the reasons below — and it does not reach significance ($p = .087$). Split by outcome, the affirmances that dominate the pooled figure show nothing at all (−1.2 points, $p = .280$, on 417 pairs). The difference lives in the 135 pairs where the appeal disturbed the judgment: −7.1 points, $p = .033$. A circuit that overturns reaches for substantive law more than the judgment it overturned did.

Two things follow, and the first is a withdrawal. On the unit at which the claim is actually made, appellate reasons are not more procedural than the reasons below; the pooled statement was an artefact of weighting pairs by their citation counts, and it is withdrawn rather than qualified. The second is that the surviving difference is reported with the same discipline applied to it: three tests appear in Table 2, and $p = .033$ would not survive correction for three any more than a model term at that level survives correction for eleven. It is reported as a signal that the strata differ, not as an established difference.

5.4 What predicts disturbance

Eleven coefficients are tested at once, so the table carries a Holm-adjusted p beside the raw one, and the text below treats as established only the three terms that survive it.

The strongest coefficient in Table 3 is the procedural share of a judgment’s own citations, at 0.74 (Holm $p = .006$). A judgment whose statutory authority is procedural is markedly harder to disturb than one resting on the law of obligations, controlling for whether it decided the merits at all. The other two survivors are the Hijri year, which is reported only inside the mature window for the reason given in §6, and one city: Dammam’s judgments are disturbed more often than the other cities’ at 1.44 (Holm $p =$

.029), which this article records and does not explain.

Three terms are worth as much as the coefficients because they are flat on every reading of the table. The number of statutory citations in a judgment (0.98), the number of distinct instruments it cites (1.05), and whether it wrote its own reasons under the headings (0.85) do not predict whether it survives. A heavily cited judgment is not a safer judgment.

Two terms sit between the two groups and are reported rather than relied on. Deciding the merits is protective at 0.86 and length is mildly aggravating at 1.11; each is significant on the raw p (.020 and .044) and neither survives correction (.157 and .307). Which way length points is worth recording even so, because it is not the direction a reader would guess: longer judgments are disturbed slightly more often, not less. Nothing in this model supports the idea that a judgment protects itself by being long.

6 Threats to validity

Only appealed judgments are here. A judgment enters this analysis because someone challenged it. Nothing here describes the judgments no one appealed, which are the great majority, and the disturbance rate is a rate among appeals rather than an error rate among judgments.

Recent judgments are censored, which is why the window is an argument. A judgment appears with an appellate decision only if the appeal was brought and decided before the corpus was collected. For the most recent years that selects the fastest appeals. The measured disturbance rate rises through the mature years and then collapses in the last two, which is a fact about the collection date and not about the courts. The model is therefore reported on 1439–1444, and the whole span is kept as the comparison that shows why.

Who writes reasons is not random, and the comparison inherits it. The reasoning comparison uses the 22.3 per cent of pairs where both levels reasoned, and reasoning goes with disturbing: a circuit that disturbs a judgment writes its own reasons far more often than one that affirms. Every statement about appellate reasoning in this article is therefore a statement about that minority of appeals, selected by their outcome, and not about appeals in general.

Eleven terms, three strata, one corpus. Testing many things on one body of data produces small p -values on its own. Both the model and the paired comparison are reported with that in mind: Table 3 carries the Holm-adjusted p and only three of its eleven terms are relied on, and the difference in Table 2 is reported as a signal rather than an established result for the same reason. Where an earlier version of this analysis pooled citations instead of pairs, the correction reversed a sign, and the withdrawal is recorded in §5 rather than removed from it.

The corpus is what the Ministry publishes. It is overwhelmingly commercial and urban. Cases that settle produce no judgment at all, and published and unpublished decisions differ. Every rate here is a rate within the published record.

The classifier is lexical. The disposition is read from operative language, validated by hand, and left unclear rather than guessed where the language does not decide it; 9.42 per cent of appellate decisions are unclear and are excluded from the rates that depend on the disposition.

7 Discussion

The finding with the widest reach is the one that requires no model. In about three appellate decisions in four, the second instance adds no reasoning of its own to the public record. That is lawful, it is efficient, and on the reason-giving literature's own account it is defensible: reasons adopted are reasons given. But it has a consequence for anyone outside the case. A practitioner who wants to know what an appellate court thinks of an argument will find, in three files out of four, only that it agreed with someone else; and a researcher measuring appellate reasoning is measuring a minority of appeals selected by their outcome. A court system that publishes its judgments has decided that its reasoning is a public good. The rate at which the second instance actually supplies it is therefore a measurable feature of that policy, and it can be reported alongside clearance rates and disposal times without anyone having to read a judgment.

The model's central coefficient joins this article to the question its predecessor asked. Measuring what Saudi courts cite shows an applied law that is overwhelmingly procedural. Measuring what survives appeal shows why that may be self-reinforcing: the procedural share of a judgment's own citations is the strongest protection against disturbance in this model. A judgment grounded in the law of procedure is harder to overturn than one grounded in the law of obligations. Whether that is because procedural questions have determinate answers, or because appellate circuits are readier to substitute their own view on substance, this data cannot say — but the incentive it describes points one way.

The nulls deserve their own sentence. Citation count, the breadth of authority cited, and whether the court below wrote its own reasons do not predict survival, and length predicts it only weakly and in the wrong direction for any account on which thoroughness protects. If judicial training or performance measurement rewards the appearance of thoroughness — a long judgment, a wall of citations — it is rewarding something this record does not connect to outcomes.

8 Conclusion

Saudi commercial judgments are published with the appellate decision inside the record of the judgment it reviews. That accident of storage yields 13,924 matched pairs and answers questions dockets cannot: not only how often appeals succeed, but whether the second instance explains itself, and what distinguishes the judgments it lets stand.

It explains itself in about one appeal in four, and more often when it disturbs than when it affirms. 10.8 per cent of appeals that reach the merits disturb the judgment. And the feature of a judgment that best protects it is not its length or the weight of authority it cites, but that its authority is procedural.

Availability of data and code

The corpus, the extraction and classification code, the model, and the scripts that generate every number in this article are public.¹ Judgments are reproduced as the Ministry publishes them, with parties' names masked at source and a further redaction pass over identifiers the publisher's masking missed.

Use of AI tools

The analysis code and drafts of this manuscript were produced with the assistance of a large language model (Anthropic's Claude) working under the author's direction. The author designed the study, framed every question, verified each result against the corpus, and is responsible for the content. No measurement in this article is produced by a language model: every number is computed by the deposited code from the deposited data.

Author biography

Abdullah Almohammed is an independent researcher in Saudi Arabia working on machine-readable corpora of Saudi legislation and adjudication. He built the corpus this article measures and the citator that joins its two halves at the level of the article. His work is published openly, with the code that produces every reported figure.

Competing interests

The author declares no competing interests. This work received no funding, and the author holds no position in, and no relationship with, the Ministry of Justice or any court whose judgments it measures.

References

- Cohen, Mathilde. 2015. "When Judges Have Reasons Not to Give Reasons: A Comparative Law Approach." *Washington and Lee Law Review* 72(2): 483.
- Eisenberg, Theodore. 2004. "Appeal Rates and Outcomes in Tried and Nontried Cases: Further Exploration of Anti-Plaintiff Appellate Outcomes." *Journal of Empirical Legal Studies* 1(3): 659–688. <https://doi.org/10.1111/j.1740-1461.2004.00019.x>
- Priest, George L., and Benjamin Klein. 1984. "The Selection of Disputes for Litigation." *Journal of Legal Studies* 13(1): 1–55. <https://doi.org/10.1086/467732>
- Shavell, Steven. 1995. "The Appeals Process as a Means of Error Correction." *Journal of Legal Studies* 24(2): 379–426. <https://doi.org/10.1086/467963>
- Siegelman, Peter, and John J. Donohue III. 1990. "Studying the Iceberg from Its Tip: A Comparison of Published and Unpublished Employment Discrimination Cases." *Law & Society Review* 24(5): 1133–1170. <https://doi.org/10.2307/3053664>

¹Repository: <https://github.com/Abdullah-m7/saudi-legal-corpus-ai>. Every measurement here is generated into a file of macros by `make_numbers.py` and typeset from it; a check refuses to build a manuscript that types one by hand.

Vogel, Frank E. 2000. *Islamic Law and Legal System: Studies of Saudi Arabia*. Leiden: Brill.