

Declared and Unapplied: What a Well-Run Legal Publisher Still Cannot Tell You About Its Own Text

Abdullah Almohammed
Independent Researcher
abdullah.m.almohammed@gmail.com
ORCID: 0009-0001-0832-0995

1 Introduction

A companion study measured whether one state’s official legal record was available and internally consistent when someone tried to build on it, and found that neither could be assumed.¹ It ended with a proposal — a five-field provenance schema that a publisher of official data could populate at the point of collection — and with a weakness it stated plainly: one jurisdiction, and an argument rather than a demonstration that the method travels.

This article is the demonstration, and it is not a replication. Repeating a measurement in a second place tests whether a number recurs. Applying an instrument in a second place tests the instrument. The second is the more useful thing to do with a proposal, and it is what this article reports: what the schema captured, what it missed, and what had to change in it.

The jurisdiction was chosen so that the answer would be informative whichever way it came out. The United Kingdom’s official legislation service is, by any ordinary standard, well run. It publishes a documented interface, permanent identifiers for every provision, and machine-readable metadata alongside every text. If the companion study’s findings were about a particular state’s capacity, almost nothing here should be measurable.

Something is, and it is not what the companion study looked for. The service publishes, in each Act’s own metadata, a list of amendments that have been enacted and that it has *not* incorporated into the text it displays. It declares this rather than concealing it. It states a target for clearing them. And a reader still cannot tell, for any particular provision, whether the declared defect is a week old or a decade old — because the field that would answer that is not published.

That is the shape of the finding, and it is the opposite of the companion study’s. There, the defect had to be excavated: four instruments whose portal text disagreed with the portal’s own amendment log, each found by hand. Here it is handed to the reader, at scale, per provision, in a form a machine can read — and it is still not enough. The gap between those two cases is what this article is about.

2 What the service publishes about its own currency

Two features of `legislation.gov.uk` matter here, and both are the publisher’s own statements rather than this study’s inferences.

The first is the *unapplied effect*. When an amendment has been made to an Act and the revised text has not yet been changed to reflect it, the service records the fact in the Act’s metadata and flags it to readers: ‘If the primary legislation on `legislation.gov.uk` has any unapplied effects, we flag them in the “Changes to Legislation” banner on the website. Clicking on the banner reveals the outstanding changes.’²

¹Cited as a companion study so this article can be read on its own; the corpus and analysis are openly published.

²*Understanding legislation*, <https://www.legislation.gov.uk/understanding-legislation>, under the heading ‘Our editorial practice and timescales’. The heading matters: the same page carries a separate section on legislation originating from the EU, and the commitment quoted here sits outside it and is stated generally.

Each such effect names the amending instrument, the provision affected, and the kind of change.

The second is a service standard: ‘We aim to incorporate new amendments into the text of the legislation within three months of those amendments coming into force.’³

The standard matters as much for what it rules out as for what it establishes. A three-month target means the expected steady state is not zero unapplied effects; a queue is the design, not the defect. Any finding here has to be about the queue’s tail rather than its existence, and the analysis is built accordingly.

3 Method

Collection. Every Public General Act from 2001 to 2018 was enumerated from the service’s own year feeds, and each Act’s metadata retrieved from its `introduction/data.xml` endpoint. That endpoint carries the complete unapplied-effects block at roughly a ninth of the bytes of the full text, which was verified rather than assumed. The service’s `robots.txt` declares a five-second crawl delay and the collector honours it, which is why a sweep of the statute book takes hours.

The schema, populated at collection time. The companion study’s five fields were written for each record as it was retrieved, unchanged from that paper’s definition. This is the point of the exercise rather than decoration: that paper claims the schema costs nothing because a collector that knows how it obtained a document already knows all five values. Populating them prospectively is the only way to find out whether that is true.

What is measured, and what is not. The companion study measured two properties: availability and internal consistency. Only the second is compared here, and the reason is structural rather than practical. Four of the schema’s five fields describe the *encounter* — whether the live source answered this collector, on this network, on this day — and encounters do not compare across collections, because they differ before the records do. No control corrects for that. Availability figures are recorded and published for this collection, and are not set against the companion study’s.

An unapplied effect has no such problem. It is a property of the data as published: any reader, on any network, at any time, obtains the same count.

4 Results

5 What the schema learned about itself

Two changes to the companion study’s proposal came out of using it, and neither would have come out of thinking about it.

A discrepancy has to record who declared it. Applied to a well-run interface, four of the five fields take a single value across every record in the collection: one source class, one retrieval route, one corroboration count, one transformation. That is the schema working, not failing — it costs nothing here because there is nothing to say.

But the defect this study measures was invisible to all five fields. An Act whose displayed text omits hundreds of enacted amendments was retrieved perfectly: official primary source, live, no transformation, nothing to report. Under the schema as written, it is indistinguishable from an Act with no unapplied effects at all.

The remedy is not a sixth field. The purpose the companion study gives `discrepancy` is ‘is there a known problem with this record?’, and a publisher flagging its own text as not current is exactly that. What the field’s *values* omit is that a discrepancy can be declared by the source as well as found by the

³Ibid., same heading.

collector — and a null means different things in the two cases. Found by a collector, it cost two or more sources and a comparison, so a null may only mean nobody looked hard enough. Declared by a source, it costs nothing to record, so a null means the publisher did not say. A single null standing for both is the collapse that paper objects to when one confidence score stands for availability and consistency at once.

The two jurisdictions are the argument. Here, 207 records carry a source-declared discrepancy and 0 a collector-found one. In the companion corpus it was the reverse: four self-contradicting instruments, every one of them found by hand. The same class of defect, with opposite provenance — and a straightforward consequence for who should be doing the work.

The fields divide by what they describe, not by what they are about. `source_class` describes the record: whose copy this is. `retrieval_route`, `corroboration` and `transformation` describe the encounter: whether the live source answered *us*, how many sources *we* found agreeing, what *we* had to do to the bytes we got. `discrepancy` is either, depending on who declared it.

Four of five describe the encounter, which is why the availability half of this comparison could not be salvaged. It also complicates the companion study's own recommendation. That paper asks publishers to populate these fields; if a publisher does, the encounter fields describe the publisher's encounter with its own source, which is a third thing again and comparable with neither. A provenance schema intended for use across collections has to mark which of its fields are encounter-relative. That one does not.

6 What follows

7 Limits

The age of an amendment is not published. Every age reported here is the age of the *amending instrument*, which is the only date the service exposes. It is not the date the amendment came into force, and the two can be decades apart: 20 effects in this collection have an affecting instrument older than the Act it amends, because provisions that were never commenced bite, when commenced, on a later consolidating Act. Compliance with the three-month target therefore cannot be computed from what the service publishes — which is itself among the findings, but it also bounds every age figure here.

Denominators are choices, and all of them are reported. An Act the service serves as enacted rather than in revised form has no revised text for an amendment to be missing from; a wholly repealed Act has no live text at all. Excluding them raises the affected share and including them lowers it. Neither is obviously right, so the article reports the share on all three bases rather than choosing one.

Coverage. 595 Acts were listed and 590 retrieved. Retrieval failures were recorded rather than retried away while availability was in scope; once it left the comparison, a second pass recovered what it could, and both attempts are kept in the record.

Primary legislation only. Statutory instruments are not covered. They are the larger population and carry their own unapplied effects; nothing here should be read as a statement about them.

Declarations

Financial support. This research received no specific grant from any funding agency, commercial or not-for-profit sectors.

Competing interests. The author declares none.

Ethics approval. Not applicable. The study involves no human participants, no animal subjects and no personal data; it analyses published national legislation and its publisher’s own metadata.

Declaration of generative AI use. *Tool:* Claude Opus 5 (Anthropic). *How it was used:* to write the collection, analysis, verification and figure scripts that produce every number reported here, to draft and revise the manuscript, and to run adversarial checks against its own claims. *Reason for use:* the study required retrieving and parsing the metadata of the entire modern statute book under a five-second crawl delay and auditing the result, which is not feasible by hand at that scale for a single researcher. The author set the research question, directed the analysis, verified the outputs against the underlying records, and is responsible for the content.

CRedit roles. A. Almohammed: conceptualisation, data curation, methodology, software, formal analysis, investigation, writing — original draft, writing — review and editing.

Data availability. The collection, the analysis and the scripts that reproduce every number and figure are openly available at <https://github.com/Abdullah-m7/saudi-legal-corpus-ai>. The underlying metadata belongs to `legislation.gov.uk` and is published by the Crown under the Open Government Licence.

A note on the subject. This article measures a defect that `legislation.gov.uk` discloses about itself. It is not an accusation. Almost no comparable publisher discloses as much, which is what makes the measurement possible at all, and the article’s argument is that disclosure of this quality still leaves a reader unable to answer the question it raises.