

Abdullah Almohammed

Independent Researcher

abdullah.m.almohammed@gmail.com

ORCID: 0009-0001-0832-0995

20 August 2026

The Editors

Language Resources and Evaluation

Springer Nature

Dear Editors,

I am pleased to submit the enclosed manuscript, “The Saudi Legal Corpus for AI: An Auditable, Official-Source-Grounded, Article-Level Corpus of Saudi Arabian Legislation,” for consideration as a Research article in *Language Resources and Evaluation*.

The manuscript describes a new language resource: a corpus of 290 Saudi legislative instruments structured into 15,689 article-level records (approximately 1.2 million Arabic tokens) across twelve legal domains, together with derived analytical layers and a retrieval benchmark. To my knowledge, no openly documented, machine-readable, article-level corpus of Saudi legislation has previously been described in the literature, and Arabic is at most marginally represented in the existing legal-NLP resource landscape (Pile of Law, MultiLegalPile, LexGLUE, LEXTREME, LegalBench).

I believe the work fits the journal’s scope in three respects. First, it is a resource paper in the strict sense: the contribution is the corpus, its construction methodology, and its documentation, rather than a modelling result. Second, it addresses resource *quality* as a first-class concern — every record carries its own provenance label, every instrument is placed in a four-tier source-verification taxonomy, and known source-staleness risks are published rather than concealed, which I offer as a reusable template for building legal corpora in jurisdictions without a machine-readable official gazette. Third, it includes an evaluation component: a manually confirmed 519-query retrieval benchmark with two deterministic baselines reported per query category, which localizes where retrieval over Arabic statutory text is actually hard (definitional queries, where the corpus’s engineered retrieval metadata outperforms BM25 by 16.4 points of top-1 accuracy).

The corpus, all derived layers, all validators, the evaluation pack, and the scripts that reproduce every figure in the manuscript are openly available under the MIT licence and archived on Zenodo with a citable DOI (10.5281/zenodo.22019183), so every statistic reported in the paper can be independently regenerated from the archived version.

The manuscript is original, has not been published previously, and is not under consideration by any other journal or conference. There are no competing interests and no funding to declare. I am the sole author. The corpus contains enacted national legislation and no personal data; the manuscript states explicitly that it is not an official government publication, contains no official translation, and does not constitute legal advice.

Thank you for your consideration.

Yours sincerely,

Abdullah Almohammed