

How Much Do Statutes Disagree?

Measuring Definitional Fragmentation Across a Legal System

Abdullah Almohammed
Independent Researcher
abdullah.m.almohammed@gmail.com
ORCID: 0009-0001-0832-0995

Abstract

Drafting manuals agree that a term should carry a stable meaning across a legal system. Whether real systems comply has never been measured, because settling it requires every statutory definition in a jurisdiction, extracted from the instrument that made it. This article supplies that measurement for Saudi Arabia — 290 instruments, 1,920 defined terms, 3,318 definitions — and argues that the measurement is worth less than the two distinctions needed to perform it. The first separates *indexical* definitions, where an instrument defines ‘the Law’ as itself, from substantive ones; indexical definitions are the most repeated and cannot disagree. The second separates lexical from substantive divergence: 61.5 per cent of shared substantive terms are worded differently, and most of that is paraphrase. Reading the 12 most widely shared terms yields 4 genuine conflicts. Fragmentation is real but narrow, and therefore a tractable target for drafting reform.

Keywords: statutory definitions; legislative drafting; interpretation; Saudi Arabia; legal corpora; terminological consistency

1 Introduction

Legislative drafting has a rule about definitions, and no evidence about how far it is kept. The rule is that a term should carry a stable meaning across a legal system, that definitions should not multiply without reason, and that a drafter departing from an established definition should know that he is departing from it. Whether legal systems comply is an empirical question that has been argued entirely from examples, because nobody has counted.

Here is one of the examples. Consider a woman who sells handmade furniture from a workshop she owns and runs alone. She has no employees. Whether she is an *establishment* — *al-munsha’ah*, a term that recurs throughout Saudi regulatory legislation — depends entirely on which statute is asking. The Competition Law defines an establishment as ‘any natural or legal person carrying on an economic activity’,¹ which she plainly is. The Labour Law defines it as ‘any undertaking run by a natural or legal person that employs one or more workers for wages, of whatever kind’,² which she plainly is not. Neither definition is defective. Each serves its own statute well. Together they mean that a single word, in a single legal system, sorts the same person onto opposite sides of a line.

¹Competition Law, Royal Decree No. M/75 (29/6/1440H) (Saudi Arabia), art. 1.

²Labour Law, Royal Decree No. M/51 (23/8/1426H) (Saudi Arabia), art. 2.

Practitioners meet this constantly. What the guidance has never had is evidence: a striking pair of divergent definitions is produced, and the reader is invited to infer either that the system is riddled with such pairs or that this one is an outlier. Both inferences are unsupported. This article counts, and the answer is neither of them.

Counting requires something few jurisdictions possess: a machine-readable corpus in which every statutory definition has been extracted from the instrument that made it, with its provenance intact. This article uses such a corpus for Saudi Arabia — 290 legislative instruments, of which 214 have a definitions article that parses cleanly, yielding 1,920 distinct defined terms carrying 3,318 definitions.³ Saudi Arabia is a useful test case for reasons beyond data availability. Its legislation was substantially rewritten during the 2020s, so definitions drafted in the 1420sH sit alongside definitions drafted last year; and its drafting convention places a definitions article at the head of nearly every instrument, which makes the definitional lexicon unusually legible.

The article's central finding is negative, and deliberately so. Measured the obvious way — comparing the wording of definitions that share a term — Saudi legislation looks badly fragmented: of the 322 substantive terms that more than one instrument defines, 198 (61.5 per cent) are worded divergently. Read by hand, that number substantially dissolves. Most of the divergence is paraphrase around a stable meaning, or homonymy, or a referent that each instrument fixes for itself and could not have fixed generally. What survives is four terms.

Getting from the first number to the second requires two separations, and those separations, rather than either number, are what this article offers to drafters and to anyone attempting a similar exercise elsewhere.

The first separates *indexical* from *substantive* definitions. Saudi instruments open by defining 'the Law' as themselves, 'the Regulation' as their own implementing regulation, 'the Minister' as whichever minister supervises them. These are by a wide margin the most repeated definitions in the corpus — 'the Law' is a defined term in 125 instruments — and they are local by construction. Two instruments defining 'the Minister' differently are not in disagreement; they have different ministers. A measure that counts them reports a drafting convention and calls it incoherence.

The second separates *lexical* from *substantive* divergence. Automatic comparison measures words; legal consistency is a property of meaning. 'Person' is defined by 27 Saudi instruments in formulations ranging from 'the natural or legal person' to 'any person of natural or legal capacity, public or private'. Some pairs share almost no vocabulary. None of them differ in law.

Section 2 sets out how Saudi statutes handle definitions and why the question is worth asking of this jurisdiction. Section 3 describes the corpus and the three filters. Section 4 gives the shape of the definitional lexicon. Section 5 examines the four conflicting terms in detail — this is the article's substantive core — and section 6 the eight that do not conflict, including the classifications that were close calls. Sections 7 and 8 draw out the consequences for drafting practice and for the legal technology now being built over statutory text, and section 9 sets out what the method cannot see.

³The corpus, its glossary layer, and the scripts that reproduce every figure in this article are openly licensed; see the data availability statement below.

2 Definitions in Saudi legislative practice

2.1 The definitions article as a convention

Saudi instruments overwhelmingly open with an article of a standard form: ‘the following words and phrases, wherever they appear in this Law, shall have the meanings assigned to them, unless the context requires otherwise’.⁴ There follows a list, sometimes of five terms, sometimes of fifty. In the corpus studied here, 214 of 290 instruments carry such an article in a form the parser could read, and the great majority of the remainder are short instruments, annexes, or organisational statutes that have no occasion for one.

Three features of the convention matter for what follows. First, definitions are *internal*: the formula binds the reader ‘wherever they appear in this Law’, and nothing in the instrument purports to bind any other. Second, definitions are *comprehensive* in aspiration — an instrument that uses a term of art usually defines it rather than relying on its ordinary meaning or on another statute’s treatment. Third, definitions are *repeated*: where an implementing regulation elaborates a law, the regulation typically restates the law’s definitions rather than incorporating them by reference, often in wording that is close but not identical.

The third feature is the engine of everything measured below. Restatement is an opportunity for drift. A drafter copying a definition from the parent instrument may modernise a phrase, drop a qualification that seems redundant, or add one that the sector requires. Each such change is individually defensible. Their aggregate is a legal system in which the same word has many masters.

2.2 No general instrument of interpretation

Common-law systems mitigate this with a general interpretation statute. The United Kingdom’s Interpretation Act 1978, Australia’s Acts Interpretation Act 1901 and Canada’s Interpretation Act each supply default meanings for recurring terms and default rules for their application, so that a drafter who says nothing about ‘person’ or ‘writing’ has nonetheless said something definite.⁵ The effect is to move a body of definitional content out of individual statutes and into one instrument where it can be maintained — and where a change to it is a legislative act with a visible author, rather than a drift that nobody decided on.

No instrument in the corpus performs that function for Saudi legislation.⁶ The consequence is structural rather than accidental: with no default layer, every instrument that needs a term must define it, and the lexicon grows by repetition. It is precisely this arrangement that makes the Saudi case worth measuring. A jurisdiction with a general interpretation act has already removed its most heavily shared terms from circulation; whatever fragmentation remains is in the tail. Saudi Arabia offers the untreated case, where the most heavily shared terms are exactly the ones still being defined instrument by instrument.

⁴The Arabic formula varies slightly between instruments; the closing qualification (*ma lam yaqtadi siyaq al-nass khilaf dhalik*) is not universal.

⁵Interpretation Act 1978 (UK), c. 30, s. 5, sch. 1; Acts Interpretation Act 1901 (Cth), pt 2 div 2; Interpretation Act (Canada), RSC 1985, c. I-21, s. 35.

⁶The Basic Law of Governance, Royal Order No. A/90 (27/8/1412H) (Saudi Arabia), is constitutional in character and supplies no general definitional apparatus. This is a statement about the 290 instruments in the corpus, which is not an exhaustive collection of Saudi legislation; see section 9.

2.3 What the literature has and has not established

Definitions are a settled topic in the drafting literature, and its received positions are largely uncontested: define only what needs defining; do not give a term a sense its ordinary meaning will not bear; and do not legislate in a definition, because a reader looking for the rule will not look there.⁷ Dickerson’s warning against the ‘humpty-dumpty’ definition — the definition that assigns a word a meaning its ordinary sense cannot bear, so that the statute says one thing and means another — remains the most cited passage on the subject. Xanthaki’s later work reframes the drafting question around effectiveness, and definitions accordingly become a question of whether the reader can find and apply the rule rather than of whether the definition is elegant.⁸

What this literature addresses is the definition *within* an instrument. The question this article asks — what happens when many instruments define the same term — sits at the edge of it, and is approached from the other side by statutory interpretation scholarship. The presumption of consistent expression, that a legislature means the same thing by the same word, is a settled interpretive canon, and so is its limit: the presumption is weak between statutes, and weaker still where the statutes serve different purposes.⁹ That limit is doctrinally sound and it is precisely what makes the empirical question interesting: if a court will not presume that ‘establishment’ means the same thing in the Competition Law and the Labour Law, then the burden of knowing that it does not falls on the reader, and nobody has told the reader how often that burden arises.

Neither literature contains a measurement. That is not an oversight; the measurement requires the definitional content of a whole legal system in a form a machine can read, and that has not been available. The legal informatics field has built large legislative corpora — Pile of Law and MultiLegalPile are the best known — but these are assembled to train language models, hold running text rather than structure, and cover no Arab jurisdiction.¹⁰ Structural standards such as Akoma Ntoso model the architecture of a legislative document faithfully but do not extract its definitional content.¹¹ Quantitative studies of legal systems have concentrated on citation structure and on document complexity rather than on terminology.¹² A companion study analyses the citation network of the same Saudi corpus; the present article takes up the complementary question of whether the instruments in that network use a common vocabulary.¹³

⁷R. Dickerson, *The Fundamentals of Legal Drafting* (Boston: Little, Brown, 2d ed. 1986), ch. 7; H. Xanthaki, *Thornton’s Legislative Drafting* (Bloomsbury, 6th ed. 2022), ch. 4.

⁸H. Xanthaki, *Drafting Legislation: Art and Technology of Rules for Regulation* (2014).

⁹Bennion, *Bailey and Norbury on Statutory Interpretation* (LexisNexis, 8th ed. 2020), ss. 21.3, 24.3; A. Scalia & B.A. Garner, *Reading Law: The Interpretation of Legal Texts* (St. Paul, MN: Thomson/West, 2012), 170–73.

¹⁰P. Henderson et al., ‘Pile of Law: Learning Responsible Data Filtering from the Law and a 256GB Open-Source Legal Dataset’, in *Advances in Neural Information Processing Systems 35, Datasets and Benchmarks Track* (2022); J. Niklaus et al., ‘MultiLegalPile: A 689GB Multilingual Legal Corpus’, in *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics* (2024).

¹¹M. Palmirani & F. Vitali, ‘Akoma-Ntoso for Legal Documents’, in *Legislative XML for the Semantic Web* (2011) 75.

¹²D.M. Katz & M.J. Bommarito, ‘Measuring the Complexity of the Law: The United States Code’, *Artificial Intelligence and Law* 22(4) (2014): 337–374, doi 10.1007/s10506-014-9160-8; C. Coupette et al., ‘Measuring Law Over Time: A Network Analytical Framework with an Application to Statutes and Regulations in the United States and Germany’, *Frontiers in Physics* 9 (2021): art. 658463, doi 10.3389/fphy.2021.658463.

¹³On the citation network of this corpus, see the companion study referred to in the data availability statement.

3 Corpus, glossary and method

3.1 The corpus and its glossary layer

The corpus comprises 290 Saudi legislative instruments — primary laws, implementing regulations, annexes and organisational statutes — captured from official and officially-sourced texts and recorded at article level. A glossary layer parses each instrument’s definitions article into term–definition pairs, retaining for each pair the term as written, the definition text as written, the instrument, and the article it came from. The layer yields 1,920 distinct terms carrying 3,318 definitions from 214 instruments.

The parse is automatic and its output is not curated. This matters for how its numbers should be read, and section 9 sets out the consequences; two concrete illustrations are worth giving here, because they are visible in the data used below. In the Unified Model Work Organization Regulation, the recorded definition of *al-munsha’ah* is a row of dots — the model form leaves a blank for the employer to complete, and the parser faithfully recorded the blank. In the Recruitment and Labour-Services Rules, the recorded definition of ‘employer’ runs on into the two definitions that follow it, because the source text’s punctuation did not mark the boundary. Neither error changes any finding reported here, but both are the ordinary condition of automatically extracted legal text, and a study that presents such extraction as clean is misleading its readers.

3.2 Filter one: indexical definitions

A definition is *indexical* when it names a particular instrument or organ of state rather than describing a class. ‘The Law’ defined as ‘the Agriculture Law’, ‘the Ministry’ defined as ‘the Ministry of Environment, Water and Agriculture’, ‘the Authority’ defined as whichever regulator the instrument creates: each of these fixes a referent for one instrument and says nothing that another instrument could contradict.

Detection uses a rule kept deliberately simple, so that any reader can inspect and disagree with it. A definition counts as indexical when it runs to at most twelve content tokens, is headed by an instrument or state-organ word (law, regulation, ministry, minister, authority, council, court, body, committee, fund, bank, centre, and so on), and does *not* open with a genus word. The last condition does most of the work. Arabic statutory drafting puts the genus first: a class definition begins ‘any natural person who...’ or ‘a document issued by...’, while a naming definition does not. A term is treated as indexical when the majority of its definitions are.

Two implementation details deserve recording, because they are the kind of thing that silently corrupts a study of this sort and neither is visible in the output.

The first is Arabic normalisation. The pipeline normalises text before comparison — stripping diacritics and *tatweel*, folding the *hamza* forms of *alif* to bare *alif*, *teh marbuta* to *heh*, *alif maqsurah* to *yeh* — as any Arabic text pipeline must. The keyword lists above, however, were written in ordinary orthography. Comparing unnormalised keywords against normalised tokens matches nothing for every head ending in *teh marbuta*, which meant that ‘the Regulation’ (*al-la’ihah*), ‘the Ministry’ (*al-wizarah*) and ‘the Authority’ (*al-hay’ah*) were classified as substantive terms, and therefore counted as being in widespread conflict with themselves. The error is invisible in the results — it produces a plausible number — and was found only by reading the classification of

individual terms.

The second is the scope of the genus test. The Arabic relative pronouns *man* and *ma* (‘whoever’, ‘whatever’) look like genus words and are common *inside* naming definitions: ‘the Authority: the body to which the Law assigns. . .’. Treated as genus words anywhere in the string, they excluded genuine indexical definitions. The test now applies only to the first two tokens, which is where the genus sits in this drafting register.

3.3 Filter two: lexical divergence

For each substantive term defined by more than one instrument, one definition per instrument is taken and pairwise Jaccard similarity computed over normalised content tokens. A term is *lexically identical* when its worst pair exceeds 0.95, *near-identical* between 0.75 and 0.95, and *lexically divergent* below 0.75. Taking the worst pair rather than the mean is deliberate: a term is only as consistent as its least consistent pair of definitions.

This measure is reported and then argued against. It is included because it is what an automatic pipeline produces and what a reader would reasonably expect such a study to report; the distance between it and the hand-adjudicated result is the point.

3.4 Filter three: hand adjudication

The twelve most widely shared substantive terms were read across every instrument that defines them and placed in one of five classes:

- **harmonised** — same concept, same scope, varied wording;
- **instrument-local** — same concept, referent fixed by the enacting instrument (each sector’s own licence, its own service provider);
- **indexical** — a naming definition the automatic filter did not catch, reported rather than silently corrected;
- **homonymous** — different concepts sharing a label;
- **conflicting** — same concept, materially different scope.

Twelve terms is a small sample, and it was chosen by frequency: these are the terms that the largest number of instruments define, and therefore the ones where inconsistency would do the most damage. The choice has a cost, taken up in section 9. No claim is made about the remaining 310 shared substantive terms beyond the lexical measure. The adjudication table is published with the analysis code so that every assignment can be checked and contested.

4 The shape of the definitional lexicon

Table 1 and Figure 1 give the shape of the data. Three features deserve comment before any interpretation.

Measure	Value
Distinct defined terms	1,920
Definition occurrences	3,318
Instruments with a parsed definitions article	214
Indexical terms	171
Substantive terms	1,749
Terms defined by more than one instrument	362
of which indexical	40
of which substantive	322
Substantive shared terms, lexically identical	94
Substantive shared terms, near-identical	30
Substantive shared terms, lexically divergent	198 (61.5%)

Table 1: Definitional structure of the corpus.




Figure 1: From all defined terms to the lexically divergent residue. Each step is a filter: 1,920 defined terms, 1,749 of them substantive, 322 of those shared by more than one instrument, 198 of which diverge lexically.

First, **most defined terms are not shared at all**. Of 1,920 terms, 1,558 — 81.1 per cent — are defined by exactly one instrument. Saudi drafting is overwhelmingly local: each statute builds its own vocabulary, and fragmentation, whatever its extent, can only live in the remaining fifth of the lexicon. This is worth stating plainly because the opposite impression is easy to form. A reader who encounters two conflicting definitions of ‘establishment’ may conclude that the system is pervasively inconsistent. The data say that the system is pervasively *local*, which is a different condition with different remedies.

Second, **indexical definitions dominate the top of any raw ranking**. The most frequently defined terms in the corpus are ‘the Law’ (125 instruments), ‘the Regulation’ (118), ‘the Minister’ (85), ‘the Ministry’ (75), ‘the Authority’ (67) and ‘the Council’ (60). All six are indexical. A study reporting these as the most contested terms in Saudi law — and a naive pipeline would report exactly that — would be reporting the definitions article convention itself.

Third, **among substantive shared terms, lexical divergence is the norm**: 198 of 322, or 61.5




Figure 2: The twelve most widely shared substantive terms, classified by hand across every instrument that defines them. Four carry materially conflicting scope; the rest diverge in wording, referent or concept rather than in law.

per cent. This is the figure that a purely automatic study would headline. Section 5 and section 6 are an argument that it should not be.

5 Four terms that conflict

Of the twelve most widely shared substantive terms, four carry materially different scope for what is recognisably the same concept. They are set out here in order of how clearly the conflict bites.

5.1 *Al-munsha'ah*: establishment

Nine instruments define *al-munsha'ah*, and they fall into three groups that do not agree.

The **competition group** defines it by reference to economic activity alone. The Competition Law has ‘every person of natural or legal capacity carrying on an economic activity’; its implementing regulation repeats this almost verbatim.¹⁴ The Implementing Regulation of the Anti-Concealment Law goes further and makes the breadth explicit: ‘everyone practising an economic activity, including the sole proprietorship, the company, and any other legal form’.¹⁵ The drafter of that instrument evidently anticipated the argument that a sole proprietorship is not an establishment, and closed it.

The **labour group** defines it by reference to employment. The Labour Law has ‘every undertaking run by a natural or legal person that employs one or more workers for wages, of whatever kind’; the High-Risk Professions Work Organization Regulation and the OSH Service Providers

¹⁴Competition Law, supra note 1, art. 1; Implementing Regulation of the Competition Law, GAC Board Decision No. 337 (25/1/1441H), art. 1 (Saudi Arabia).

¹⁵Implementing Regulation of the Anti-Concealment Law, issued under art. 19 of the Anti-Concealment Law, Royal Decree No. M/4 (1/1/1442H), art. 1 (Saudi Arabia).

Regulation each restate this in near-identical terms.¹⁶ An entity with no employees is not an establishment for any of them.

The **social insurance group** defines it by reference to legal form and identity: ‘the legal entity through which the employer conducts his activity, and which bears his name or the trade name of that activity’.¹⁷ This is a third thing again — neither activity nor headcount but the vehicle.

The divergence is not stylistic, and its consequences are not theoretical. Our sole trader is an establishment under the competition and anti-concealment regimes, and therefore within the reach of obligations those instruments impose on establishments; she is not an establishment under the labour and occupational-safety regimes, and therefore outside obligations those impose. Whether that is the right allocation is a policy question on which this article takes no position. The point is narrower and, we think, harder to argue with: the allocation is achieved by a single word carrying three different meanings, which is a fragile way to achieve anything. A drafter extending an establishment-based obligation into a new instrument must know which of the three definitions he is inheriting, and the corpus shows that there is no default answer.

There is a fourth group, of one. The Unified Model Work Organization Regulation, an annex to the Labour implementing regulation, ‘defines’ *al-munsha’ah* as a row of dots, because it is a model form for employers to complete. It is a parsing artefact rather than a legal position, and it is recorded here rather than quietly dropped, because the count of nine instruments includes it.

5.2 *Al-mamlakah*: the Kingdom

Seventeen instruments define ‘the Kingdom’, and the division between them is unusually clean. Thirteen define it as ‘the Kingdom of Saudi Arabia’ and stop — a definition that does no work beyond confirming that the ordinary meaning is intended.

The other four define territory. The Value Added Tax Law and the Real Estate Transaction Tax Law each define the Kingdom as ‘the territory of the Kingdom of Saudi Arabia, including the areas outside the territorial waters over whose waters, seabed, subsoil and natural resources the Kingdom of Saudi Arabia exercises sovereign rights and jurisdiction under its own law and international law’.¹⁸ The Income Tax Law reaches the same result by a longer route, adding airspace and the Kingdom’s rights in the Saudi–Kuwaiti Divided Zone.¹⁹ The Zakat Collection Implementing Regulation has simply ‘the territory of the Kingdom of Saudi Arabia’, without the maritime extension.²⁰

What makes this instructive is that the four expansive definitions are not scattered: they are exactly the four fiscal instruments in the set. The extended definition is a tax-regime phenomenon, and it is easy to see why — fiscal jurisdiction over offshore activity has to be stated, because it is not obvious. But this also means the divergence is invisible to anyone reading a single instrument,

¹⁶Labour Law, *supra* note 2, art. 2; High-Risk Professions Work Organization Regulation art. 1; OSH Service Providers Licensing and Accreditation Regulation art. 1 (Saudi Arabia).

¹⁷Implementing Regulation of the Social Insurance Law, Ministerial Resolution No. 230/*ta’minat* (26/12/1445H), art. 1 (Saudi Arabia); the legacy regulation carries the same wording.

¹⁸Value Added Tax Law, Royal Decree No. M/113 (2/11/1438H), art. 1 (Saudi Arabia); Real Estate Transaction Tax Law, Royal Decree No. M/84 (19/3/1446H), art. 1 (Saudi Arabia).

¹⁹Income Tax Law, Royal Decree No. M/1 (15/1/1425H), art. 1 (Saudi Arabia).

²⁰Zakat Collection Implementing Regulation, Minister of Finance Resolution No. 1007 (19/8/1445H), art. 1 (Saudi Arabia).

and consequential where regimes meet. An offshore activity may sit inside ‘the Kingdom’ for VAT purposes and outside the definition adopted by a sectoral regulator whose instrument says only ‘the Kingdom of Saudi Arabia’. Whether the sectoral instrument’s bare definition was intended to exclude offshore areas, or whether its drafter simply never considered the question, is not something the text can tell us — which is itself the problem. A term that fixes the geographical scope of a statute is the last place a legal system should leave a silent divergence.

The zakat variant illustrates a further point. ‘The territory of the Kingdom’ is not the same as ‘the Kingdom of Saudi Arabia’ and not the same as the VAT formula; it is a third position, adopted in an instrument that sits directly alongside the other three fiscal instruments and is administered by the same authority. If definitional drift can occur within a single revenue administration, it can occur anywhere.

5.3 *Al-mustahlik*: consumer

Thirteen instruments define ‘consumer’, and unlike ‘establishment’ the divergence here is not two or three camps but a spread. The Credit Information Law has ‘the natural or legal person having credit dealings’.²¹ The Electricity Law and its regulation have ‘every person supplied with electrical service for his own consumption’; the Water Law has ‘every person who uses water for his own consumption’. The E-Commerce Law has ‘the person dealing in electronic commerce with a view to obtaining the products or services which the service provider makes available’.²² The Payment Systems and Services Law has ‘the person to whom the ultimate benefit of payment services is directed’.²³ The Product Safety Law has ‘a natural or legal person who uses a product’. The Dry Gas and LPG Distribution Law defines the consumer as ‘the residential consumer or the commercial consumer’ — a definition by enumeration that presupposes two further terms.

Two features distinguish these from one another in ways that matter. Several of the definitions are limited to *own consumption* (electricity, water); several are not (product safety, e-commerce). And several confine the term to a *relationship with a licensed supplier*, while product safety attaches it to use of a product regardless of any relationship at all.

The regimes are adjacent enough for this to bite. A household buying a gas appliance online for use with a metered supply is a consumer under the E-Commerce Law by reason of the transaction, under the Product Safety Law by reason of use, and under the Dry Gas Distribution Law only if it falls within ‘residential consumer’ as that instrument’s own further definitions provide. The protections attaching to each status differ. Whether they were intended to differ in this particular combination is doubtful; nothing suggests the question was ever put.

5.4 *Al-nashat*: activity

Ten instruments define ‘activity’, and this is the case where the classification was most contestable. The environmental family — five implementing regulations sharing near-verbatim wording — has ‘any project, or industrial, commercial or service facility, or works, or other activities expected to

²¹ Credit Information Law, Royal Decree No. M/37 (5/7/1429H), art. 1 (Saudi Arabia).

²² E-Commerce Law, Royal Decree No. M/126 (7/11/1440H), art. 1 (Saudi Arabia).

²³ Payment Systems and Services Law, Royal Decree No. M/26 (22/3/1443H), art. 1 (Saudi Arabia).

have environmental effects’.²⁴ The Income Tax Law has ‘commercial activity in all its forms, or professional or craft activity or any similar activity, intended to realise profit, including the use of movable and immovable property’. The Zakat Collection Implementing Regulation has ‘every work intended to generate income on a continuous and regular basis under a licence issued by the competent authorities’. The Contractors Classification Law has ‘a type of subsidiary work within a field’. The Railway Law has ‘an activity connected with the provision of railway services’. The tenth, the Recruitment and Labour-Services Rules, has ‘the activity the licensee is permitted to practise under these Rules’ — an instrument-local definition, and another instance of the run-on defect described in section 3.1.

The environmental definition and the fiscal definitions are doing genuinely different jobs: one identifies conduct that triggers environmental control, the other identifies conduct that generates taxable income. On one reading that is homonymy, and the term should join ‘permit’ in section 6 below. It is classified as conflicting here for a narrower reason: the environmental and fiscal regimes both attach obligations to the same commercial conduct, so an undertaking can be asked in one breath whether an activity is ‘an activity’, and the two answers can differ while the conduct does not. The classification is a judgement, it is recorded as such, and a reader who reclassifies it moves the headline from four conflicts to three.

6 Eight terms that do not conflict

The remaining eight of the twelve are as important to the argument as the four, because each is lexically divergent and none of them represents a legal inconsistency. They fall into four groups.

6.1 Harmonised: person, licensee, employer

‘Person’ is the extreme case and the most instructive. Twenty-seven instruments define it. The wording ranges from ‘the natural or legal person’ through ‘any person of natural or legal capacity, public or private’ to ‘any natural or legal person recognised as such by the laws of the Kingdom of Saudi Arabia’. Pairwise lexical overlap between some of these is close to zero. The legal content is the same in every case: natural persons and legal persons, with some instruments adding the public/private clarification that others leave implicit.

Two definitions in the set are not like the others, and both are worth stating. The Real Estate Transaction Tax Law extends the term to ‘any other entity to which statutory provisions grant the right to conclude real-estate dispositions’, which reaches arrangements that may not be legal persons in the ordinary sense. And the Human Organ Donation Law defines ‘person’ as ‘every human being of full capacity who has passed eighteen years’ — a definition that excludes legal persons entirely and adds an age threshold.

That second definition is a genuine boundary case, and it was classified as harmonised rather than conflicting for a reason that should be stated so it can be disputed. The organ-donation definition operates in a domain where the commercial sense has no application: nothing turns on

²⁴Implementing regulations made under the Environmental Law, Royal Decree No. M/165 (19/11/1441H) (Saudi Arabia): the Air Quality Regulation, the Environmental Fees Regulation, the Environmental Inspection and Audit Regulation, the Environmental Permits Regulation and the Environmental Service Providers Regulation, art. 1 in each. The Environmental Law itself does not define the term.

whether a company is a ‘person’ for the purposes of consenting to donate an organ, because a company cannot. There is no transaction in which the two definitions could be applied to the same subject and yield different answers. On a stricter view this is homonymy — the same label over two concepts — and it would then belong in section 6.4 rather than here. What it is not, on any view, is a conflict of the kind ‘establishment’ presents, where two instruments compete to classify the same real person.

‘Licensee’ (21 instruments) and ‘employer’ (8) are harmonised in a plainer way. Every definition of ‘licensee’ amounts to the holder of the licence that the instrument in question issues — with the electricity instruments extending it to the holder of ‘a licence or a valid exemption’, which widens the class without changing the concept. Every definition of ‘employer’ amounts to a person who employs one or more workers for wages, with the social insurance instruments substituting ‘worker or employee’ and ‘wage or salary’ to match their own coverage.

‘Employer’ has one outlier worth stating. The Saudization Mediation Rules define the employer as ‘every natural or legal person requesting one or more Saudi workers’ — a definition of the party to the mediation rather than of the employment relationship. It is not classified as a conflict because no obligation anywhere turns on treating such a person as an employer generally: the definition operates only inside the mediation those Rules establish. But it is a definition of a different thing under the same label, and a reader who took it out of its instrument would go wrong.

6.2 Instrument-local: licence, service provider

‘Licence’ is defined by 26 instruments and ‘service provider’ by 9, and in both cases the definitions differ because they must. A licence, in each instrument, is the document that that instrument’s regulator issues to permit that instrument’s regulated conduct: a written permission from the Ministry to practise an agricultural activity; a document from the Authority to provide postal services; an approval from the Ministry to practise one or more of the activities subject to article 5 of the Energy Supplies System. The concept is constant — an authorisation, its issuer, and the conduct it authorises — and the three variables are filled in locally, as only the enacting instrument can fill them.

This class is easy to confuse with conflict and it is the opposite of conflict. A general definition of ‘licence’ would be useless: it would have to be so abstract as to decide nothing. The variation is the definition doing its job.

6.3 Indexical definitions the filter missed

Two of the twelve — ‘relevant bodies’ (8 instruments) and ‘supervising body’ (8) — are designating definitions that the automatic filter classified as substantive. ‘Supervising body’ is defined by the Associations and Non-Profit Institutions Law as ‘the government body within whose competences the association’s or institution’s activity falls’, by the environmental instruments as ‘any government body legally authorised to supervise the activity’, and by the Implementing Regulation of the Law of the Saudi Council of Engineers as, simply, ‘the Ministry of Municipalities and Housing’. ‘Relevant bodies’ is defined by the Building Code Law by enumerating three named ministries and the Saudi Standards, Metrology and Quality Organization, and by the Child Protection Law as ‘the bodies concerned with child protection, as the Regulation shall determine’.

These name, for each instrument, whichever organs of state that instrument's subject matter engages. They cannot conflict, for the same reason 'the Minister' cannot. The filter missed them because most of their definitions are phrased as classes — 'any government body legally authorised to...' — rather than as names, which is precisely the shape the filter uses to identify a substantive definition.

The classification is itself a judgement at the margin: a reader might reasonably call these instrument-local rather than indexical, since they behave much as 'licence' does. Nothing in the results turns on the choice — both classes are non-conflicting — but it is a real boundary, and the two categories are less distinct than the five-way scheme suggests.

These misses are reported rather than corrected. Correcting them would raise the filter's apparent accuracy on the very sample used to assess it. Taken as a recall estimate on widely shared terms, ten of twelve is around 83 per cent, and that is the only recall figure this study can offer.

6.4 Homonymy: permit

'Permit' (*al-tasrih*, 11 instruments) is the one clear case of two concepts sharing a label. Seven of the eleven definitions — the environmental family, electricity, water — describe a document issued before an activity begins, in respect of preparatory work. The Waste Management Law has a compliance certificate issued to recycling facilities before they obtain their licences — close to that family but not the same thing. The Agriculture Regulation has a one-off ministerial permission for temporary agricultural operations. The Civil Aviation Law has something else entirely: 'permission for an aircraft to cross the Kingdom's airspace or land at one of its airports for commercial or diplomatic purposes'.²⁵ The Telecommunications and Information Technology Act has a third thing: a registration issued to a provider located outside the Kingdom that offers services into it.²⁶

An overflight clearance and an environmental pre-commencement document have nothing in common except the Arabic word. That is a feature of natural language, not a defect of legislation, and no harmonisation project should touch it. It is included in the twelve because a lexical measure cannot tell it apart from conflict — which is the whole difficulty.

7 What follows for drafting

The most useful thing this study produces for a drafting body is not the 61.5 per cent and not the 4 terms, but the demonstration that the second number can be obtained at all.

If definitional conflict in a legal system were pervasive, harmonisation would be a generational undertaking and probably not worth starting. If it is concentrated in a handful of high-traffic terms, it can be addressed deliberately. This study's estimate — four conflicts among the twelve most widely shared substantive terms — is not a system-wide count, and section 9 explains why it is more likely a floor than a ceiling. But the shape it suggests is a system that is definitionally *local* rather than definitionally *incoherent*, and the difference matters for what one should do about it.

Three responses suggest themselves, in ascending order of ambition.

²⁵Civil Aviation Law, Royal Decree No. M/44 (18/7/1426H), art. 1 (Saudi Arabia).

²⁶Telecommunications and Information Technology Act, Royal Decree No. M/106 (2/11/1443H), art. 1 (Saudi Arabia).

The least ambitious is a **definitions register**: a maintained list of terms defined by more than one instrument, with their definitions and the instruments that made them, available to drafters at the point of drafting. A drafter about to define ‘establishment’ would see the three existing definitions and choose among them knowingly. This requires no legislative change and no harmonisation — only that the information exist in a usable form. The corpus used here shows that it can be assembled from published texts, and the article’s method shows how to strip out the indexical noise that would otherwise make such a register useless.

The middle response is a **drafting-stage check**. The pipeline described in section 3 runs over a corpus, but nothing prevents it running over a draft: given a proposed definitions article, it can report which of the proposed terms are already defined elsewhere and how far the proposal diverges. That is a lint check for legislation, and its natural home is the body that reviews drafts before enactment. Its output would be advisory — as section 5 shows, divergence is often correct — but advisory information that a drafter currently does not have.

The most ambitious is a **general instrument of interpretation**, on the model of the Commonwealth interpretation acts, supplying default definitions for the terms that recur across the system. The corpus makes the candidate list computable: ‘person’, defined in 27 instruments with no legal variation between them, is the obvious first entry, and the exercise of drafting a default definition would surface the two outliers discussed in section 6.1 as express exceptions rather than as silent divergence. That is a substantial legislative undertaking, and this article does not argue that Saudi Arabia should embark on it. It does suggest that the question is now answerable with evidence rather than intuition, and that the evidence points to a small enough set of terms for the undertaking to be finite.

It is worth being precise about how much a general interpretation instrument would actually achieve here, because the answer is instructive. Of the four conflicts in section 5, it would resolve one cleanly and none of the others. ‘The Kingdom’ is exactly the kind of term such an instrument handles well: a single default definition, with the fiscal instruments expressly displacing it where offshore jurisdiction is intended, would convert a silent divergence into a stated one. ‘Establishment’, ‘consumer’ and ‘activity’ it would not resolve, because in each case the divergence tracks a real difference in regulatory purpose that a default definition would have to be displaced to accommodate — and a default that is displaced by every instrument that uses the term has achieved nothing. What it *would* achieve, even there, is that each displacement would be an express choice recorded in the displacing instrument. That is a smaller benefit than harmonisation and a more durable one.

A fourth response should be named in order to be set aside: harmonising the four conflicting terms themselves. That would be the wrong lesson. The divergence between the competition and labour definitions of ‘establishment’ may well be correct; the two regimes regulate different things and have different reasons to draw the line where they draw it. What is not defensible is that the divergence is undocumented, that no instrument acknowledges it, and that the reader of either statute has no way of learning that the other exists. The remedy for a known divergence is a cross-reference. The remedy for an unknown one is to find it first.

8 What follows for legal technology

Systems that answer questions over statutory text must resolve terms, and the conflicting cases are exactly where naive resolution fails.

The failure mode is specific enough to state precisely. A retrieval-augmented system asked whether a sole trader is subject to an obligation imposed on ‘establishments’ will retrieve a definition of ‘establishment’ — and which definition it retrieves determines the answer it gives. Retrieval ranks by similarity to the query, and the query in this example contains nothing that distinguishes the competition regime from the labour regime. The system will therefore answer confidently, and it will be right or wrong according to which of three definitions happened to rank first.

The lesson is that term resolution in legal text is not a lookup but a choice of governing instrument, and that the choice must be made before the definition is retrieved rather than after. A system that models it will ask which statute the question arises under; a system that does not will be confidently wrong on precisely the terms that matter most, because the terms that matter most are the ones many instruments define.

This also gives a concrete use for the kind of enumeration this article performs. A list of terms known to conflict is a list of the points at which such a system must refuse to answer without knowing the governing instrument. Four terms is a short enough list to encode as a rule.

9 Limitations

The glossary is parsed, not curated. Definitions come from an automatic parser over definitions articles. Section 3.1 gives two visible examples of what that costs: a definition recorded as a row of dots, and a definition that runs on into the two that follow it. Terms defined outside the definitions article — in a substantive provision, or in a schedule — are not captured at all. The counts here are approximate at the margins, and should be read as such.

The indexical filter is a heuristic with one recall estimate. It is inspectable, and its misses on the adjudicated sample are reported: two of 12, giving 83.3 per cent recall on widely shared terms. It has no measured precision, and no evaluation at all against a purpose-built hand-labelled sample. Building that sample is the obvious next step for anyone extending this work.

The adjudication is single-annotator and small. Twelve terms were classified by one reader. There is no inter-annotator agreement to report; the classes involve judgement at the boundaries, and sections 5.4 and 6.3 identify two places where a different reader could reasonably classify differently. The finding should be read as a documented and checkable reading, not as a measurement.

Frequency selection probably understates conflict. The twelve adjudicated terms are those that the most instruments define, and frequency selects for generality: widely shared terms tend to be broad concepts (‘person’, ‘licence’) that are easy to harmonise or that vary only by sector. Genuinely conflicting *technical* terms — shared by two or three instruments in adjacent regimes, where drafters are most likely to assume they are borrowing a settled meaning — sit in the tail and

are not examined here. Four in twelve is therefore a lower bound on the density of conflict in the shared lexicon rather than an estimate of it.

Divergence is not the same as inconsistency. Two instruments may define a term differently for good reason. Nothing here establishes that the four conflicting cases are drafting errors, and section 7 argues expressly against assuming they are. The claim is that they are places where one word carries different legal content, which is a fact about the system that its users need to know.

Coverage. Only the 214 instruments whose definitions articles parse cleanly contribute. Instruments outside the corpus, instruments without a definitions article, and instruments the parser skipped are invisible to the study. The corpus is substantial but it is not the whole of Saudi legislation, and statements above about what no instrument does — including the observation that no general interpretation instrument exists — are statements about the corpus.

10 Conclusion

Definitional fragmentation in Saudi legislation is real, narrow, and measurable, and the gap between the second and third of those adjectives is where the difficulty lies.

Measured automatically, the system looks fragmented: 61.5 per cent of shared substantive terms are worded divergently. Read by hand, most of that divergence turns out to be paraphrase, homonymy, or a referent each instrument fixes for itself. Among the twelve terms that the largest number of instruments define, four carry materially different legal scope. Within that sample, in other words, lexical divergence overstates substantive conflict roughly threefold — a ratio measured on twelve terms, which should not be projected onto the other 310 shared substantive terms and still less onto other jurisdictions. Measuring it elsewhere would be a useful test of whether anything like it holds generally.

The methodological claim is the one that travels. Any attempt to measure definitional consistency automatically will meet both of the separations set out here. Indexical definitions exist wherever instruments name themselves and their own institutions, which is to say nearly everywhere; and lexical similarity is a poor proxy for legal identity in any language. A study that performs neither separation will report a number about drafting convention and paraphrase, and will present it as a number about law.

The four conflicting terms are the smallest finding this study could have produced and, for that reason, the most useful one. *Establishment* means one thing to the competition regime and another to the labour regime, and our sole trader with her furniture workshop falls on different sides of that line depending on which statute is asking. A practitioner in either field may know this. Neither statute says it. And for the first time in this jurisdiction, it can be enumerated rather than encountered.

Acknowledgements and declarations

Funding. No funding was received for conducting this study.

Competing interests. The author declares no competing interests.

Ethics approval. Not applicable. The study involves no human participants, no animal subjects and no personal data; it analyses published national legislation only.

Data availability. The corpus, its glossary layer, and the analysis and figure scripts that reproduce every number and figure reported here are openly available at <https://github.com/Abdullah-m7/saudi-legal-corpus-ai> and archived on Zenodo under the MIT licence (10.5281/zenodo.22019183; concept DOI 10.5281/zenodo.22019182). The companion citation-network study referred to in section 2.3 is under review; a preprint is available from the same repository.

Legal status. This article analyses a non-official research corpus. Nothing here constitutes legal advice, and the binding text of any instrument is the Arabic original published in *Umm al-Qura*. Instrument identifiers are given as recorded in the corpus registry.