

The Code the Courts Do Not Cite:

Measuring Uptake After Codification

Abdullah Almohammed¹

ABSTRACT

Comparative law has long asked what happens when a state enacts a code. The answer has been sought in doctrine, in institutions, and in economic aggregates, but almost never in the one place where a code either works or does not: the reasoning of the courts said to apply it. This Article supplies that measurement for a jurisdiction that has just codified. Saudi Arabia enacted a comprehensive civil code in 1444H, and the same royal decree amended the jurisdiction clause of its commercial courts to point at it. Its Ministry of Justice publishes commercial judgments in full text. Matching 121,207 statutory citations drawn from 50,666 published judgments to 15,855 enacted articles across 291 instruments, this Article measures uptake at the level at which a lawyer argues: the article.

Uptake is low and lopsided. 11.7 per cent of the statute book is ever cited; 30.5 per cent of the articles within the commercial courts' own statutory jurisdiction; and 89.2 per cent of what is applied is procedural rather than substantive. The contrast needs no denominator: the law telling the court how to proceed has 99.0 per cent of its articles cited, and the code telling the court what the parties owe each other has 27.5 per cent. Segmenting each judgment by its own headings locates the narrowing at the bench rather than the bar. The Article argues that enactment and uptake are separate facts, that the second is measurable wherever judgments are published, and that comparative law's central claims about transplanted codes have been resting on the first.

¹Independent researcher, Kingdom of Saudi Arabia. ORCID 0009-0001-0832-0995. The corpus, the extraction code, and the scripts that generate every figure in this Article are public; see *infra* note 32.

I. Introduction

A legislature that enacts a code has done something. What it has done is usually described in the language of the enactment itself: a rule now exists, a gap is closed, a field is governed. Whether the courts said to apply the code in fact reason from it is a different proposition, and it has been remarkably hard to check.

The difficulty is not conceptual. It is that checking requires two things at once, in the same jurisdiction, at the same level of detail: a machine-readable statute book, article by article, and a body of judgments in full text large enough that what the courts cite can be counted rather than sampled. Where judgments are published in full, codes are old; where codes are new, judgments are seldom published. Comparative law has therefore argued about the effect of transplanted codes for half a century on evidence that never included the courts' own words at scale.²

Saudi Arabia has, for a brief period, both. The Civil Transactions Law was issued by Royal Decree No. M/191 of 29/11/1444H, codifying the general law of obligations and contract for the first time in the Kingdom's modern history.³ The same decree amended article 16 of the Commercial Courts Law, the provision that fixes what the commercial courts may hear.⁴ And the Ministry of Justice publishes commercial judgments in full text, including the appellate decision where a case went up.

²The canonical statement of the problem is ALAN WATSON, *LEGAL TRANSPLANTS: AN APPROACH TO COMPARATIVE LAW* (2d ed. 1993), which treats borrowing as the ordinary engine of legal change. The principal empirical challenge is Daniel Berkowitz, Katharina Pistor & Jean-François Richard, *The Transplant Effect*, 51 *AM. J. COMP. L.* 163 (2003), which measures the consequences of transplantation through institutional and economic indicators rather than through adjudication itself.

³Civil Transactions Law, Royal Decree No. M/191 (29/11/1444H) (Saudi Arabia) [hereinafter CTL]. The corpus used here holds all 721 of its articles, each verified against the Ministry of Justice's own legal portal.

⁴Commercial Courts Law, Royal Decree No. M/93 (15/8/1441H) (Saudi Arabia) [hereinafter CCL], art. 16, as amended by Royal Decree No. M/191 (29/11/1444H). The amendment is not incidental to this Article's question: the instrument that created the code also redrew, in the same breath, the jurisdiction of the courts most likely to apply it.

That conjunction permits the measurement this Article reports. Drawing 121,207 statutory citations from 50,666 published judgments and matching each to the article cited, across 15,855 articles in 291 instruments, it asks how much of an enacted statute book a court system actually uses.

The answer is that uptake is small, and that its shape is more informative than its size. 11.7 per cent of the articles in the statute book are ever cited in a published judgment. Narrowing the denominator to the articles falling within the commercial courts' own statutory jurisdiction — the field the legislature itself marked out for them — raises it to 30.5 per cent, which still leaves better than two articles in three untouched. And 89.2 per cent of everything the courts do cite is procedural: it governs how a case is heard rather than what the parties owe each other.

The sharpest form of the result needs no denominator at all, because it compares two instruments to themselves. The Commercial Courts Law, which tells the court how to proceed, has 99.0 per cent of its 96 articles cited somewhere in the corpus. The Civil Transactions Law, which tells the court what the parties owe each other — and which the amended jurisdiction clause of the first statute points the court toward — has 27.5 per cent of its 721. A legislature enacted both. The courts have taken up one.

Three further results bear on why, and each is designed to close an explanation that would otherwise be available. First, the narrowing is not imposed by the parties. Segmenting each judgment at its own headings, and attributing each citation in the statement of the case to the court or to a litigant, shows the court's reasons to be *more* procedural than the arguments put to it, not less.⁵ Substantive instruments are raised, and are answered with procedure.

Second, the unused portion of the book is not merely its neglected or superseded portion. If legislative attention tracked judicial use, the amended part of the statute book would be the applied part. Within instruments, it is not.

Third, the measurement is not an artefact of what the Ministry chooses to publish, or at least not

⁵See *infra* Part IV.C. The attribution step matters: a court reciting the provisions under which it convened is not a party raising them, and leaving that narration inside the parties' segment flatters the parties' procedural share.

in the direction that would explain it away. The published record is overwhelmingly commercial, which is precisely why the Civil Transactions Law's low share is difficult to dismiss: these are the courts whose jurisdiction clause was amended to name it.

This Article proceeds in five Parts. Part I sets out what comparative law has claimed about codification and transplantation, and why the claim has not been tested against adjudication. Part II describes the Saudi setting: the code, the jurisdiction clause, and the publication practice that makes the measurement possible. Part III states the method, including three defects that produced confident wrong answers before they were found. Part IV reports the findings. Part V takes up what follows for comparative law, for legislative evaluation, and for the courts themselves, and states the limits that bound every figure here — the most important being that this Article measures the first years of a code, and a code's first years are not its verdict.

II. Codification and Its Unmeasured Promise

A. What Enactment Is Taken to Accomplish

A code promises three things at once, and they are worth separating because only the first is delivered by the act of enactment.

The first is *availability*. Where the governing rule previously had to be derived — from doctrinal treatises, from the reasoning of earlier adjudications, from the opinions of qualified jurists — a code supplies a numbered article that can be read, quoted, and disagreed with. This is accomplished the moment the text is promulgated.

The second is *determinacy*. A code is supposed to narrow the range of defensible answers to a legal question. Two judges reading article 76 should reach the same place more often than two judges reasoning from first principles. This is not accomplished by promulgation; it is a prediction about what judges will do.

The third is *displacement*. A code is supposed to become the operative source, so that the uncoded materials it drew on recede into a background role. This too is a prediction, and it is the strongest of the three, because a code may be available and even determinate while the courts

continue to decide on other grounds entirely.

Comparative scholarship has generally treated the second and third as following from the first, or as questions to be answered doctrinally — by reading the code, reading the commentary, and reasoning about what a court applying them would do.⁶ This Article's claim is narrower than a criticism of that method: it is that where judgments *are* published in full, displacement is an observable fact rather than an inference, and that observing it changes what can be said.

B. Transplantation, Irritation, and the Missing Evidence

The question has been most sharply posed in the literature on legal transplants. Watson's position, that borrowing is the ordinary mechanism of legal change and that rules travel more easily than their social context would suggest, made transplantation look cheap.⁷ The reaction to it made transplantation look expensive in two different ways.

Legrand's objection was conceptual: a rule detached from the interpretive community that gives it meaning is not the same rule, so a transplant in the strict sense is impossible.⁸ Teubner's was mechanical: a borrowed doctrine does not integrate but irritates, and the receiving system reorganizes around it in ways neither the lender nor the borrower intended.⁹ Langer's refinement, published in this journal, replaced the metaphor of transplantation with translation, and showed through plea bargaining's spread into four civil-law systems that what arrives is systematically not what was sent.¹⁰

⁶The move is not naive. Where judgments are unpublished or terse, doctrinal inference is the only method available, and much of the comparative literature was built in exactly that condition.

⁷ALAN WATSON, *LEGAL TRANSPLANTS: AN APPROACH TO COMPARATIVE LAW* 21–30, 95–101 (2d ed. 1993).

⁸Pierre Legrand, *The Impossibility of 'Legal Transplants'*, 4 MAASTRICHT J. EUR. & COMP. L. 111 (1997).

⁹Gunther Teubner, *Legal Irritants: Good Faith in British Law or How Unifying Law Ends Up in New Divergences*, 61 MOD. L. REV. 11 (1998).

¹⁰Máximo Langer, *From Legal Transplants to Legal Translations: The Globalization of Plea Bargaining and the Americanization Thesis in Criminal Procedure*, 45 HARV. INT'L L.J. 1 (2004).

Alongside this ran an empirical literature that tried to price the transplant directly. Berkowitz, Pistor and Richard argued that how a legal order was received — whether it was adapted to local conditions and whether the population was familiar with its concepts — predicted later legal effectiveness better than which family it came from.¹¹ That work descends from the legal origins literature, which correlated the content of formal law with financial and economic outcomes across many jurisdictions.¹²

Every one of these accounts is about what happens after a code arrives. Not one of them observes the courts applying it, article by article. Legrand and Teubner argue from doctrine; Langer works from close comparative study of procedure in four systems; Berkowitz, Pistor and Richard work from country-level indices of legality; the legal origins literature works from the content of statutes and from economic aggregates. The dependent variable is always something other than the reasoning of the courts said to be applying the transplanted law.

That is not an oversight. Measuring the reasoning requires the reasoning to be public, in quantity, and machine-readable, and requires the statute book to be available at the same granularity — and until very recently no jurisdiction that had just codified offered both.

C. What Would Count as Evidence

If the operative question is displacement — whether the enacted code has become the source the courts reason from — then the evidence is what the courts cite, and the unit is the article.

The unit matters more than it may appear. An instrument-level measurement asks whether a code has been cited at all, and every enacted code will pass that test within months: one citation in one judgment suffices. An article-level measurement asks how much of the code has entered

¹¹Daniel Berkowitz, Katharina Pistor & Jean-François Richard, *The Transplant Effect*, 51 AM. J. COMP. L. 163 (2003). The economic companion is Daniel Berkowitz, Katharina Pistor & Jean-François Richard, *Economic Development, Legality, and the Transplant Effect*, 47 EUR. ECON. REV. 165 (2003).

¹²Rafael La Porta, Florencio Lopez-de-Silanes, Andrei Shleifer & Robert W. Vishny, *Law and Finance*, 106 J. POL. ECON. 1113 (1998).

use. A code of several hundred articles that is cited constantly, but always through the same dozen provisions, has been received in a way that instrument-level counting cannot distinguish from full uptake — and the difference is the whole of what codification was supposed to change.

Three conditions therefore have to hold together, and they rarely do:

1. The statute book must exist as structured text at article level, with each article's number, status and text recoverable, so that a citation can be matched to the provision it names rather than merely to the instrument.
2. Judgments must be published in full text and in quantity, so that what courts cite can be counted rather than sampled, and so that *where* in a judgment a citation falls — in the recital of the parties' arguments or in the court's own reasons — can be read.
3. The code must be new enough that uptake is still a live question. Measuring uptake of a two-hundred-year-old code answers a different question: it describes a settled equilibrium rather than the transition into one.

Work joining legislation and adjudication computationally has grown quickly, but it has grown where the first two conditions hold and the third does not.¹³ The United States and Germany are not jurisdictions in which anyone is waiting to see whether the code will be taken up.

Saudi Arabia satisfies all three, and will not satisfy the third for long. That is the opportunity this Article takes.

III. The Saudi Setting

A. A Statute Book Built in a Decade

The Kingdom's constitutional arrangement makes the question here unusually sharp. The Basic Law of Governance provides that the courts apply the rules of the Islamic Shari'ah, together with the

¹³See, e.g., Corinna Coupette et al., *Measuring Law Over Time: A Network Analytical Framework with an Application to Statutes and Regulations in the United States and Germany*, 9 FRONTIERS IN PHYSICS, art. 658463 (2021) (statutory structure over time in two mature systems); James H. Fowler et al., *Network Analysis and the Law: Measuring the Legal Importance of Precedents at the U.S. Supreme Court*, 15 POL. ANALYSIS 324 (2007) (citation networks among judicial decisions).

regulations issued by the ruler.¹⁴ A Saudi judge therefore has, as a matter of constitutional structure, a source of decision that does not depend on any enacted text. Codification in this setting is not the reorganization of an existing statutory field; it is the introduction of enacted articles alongside a source that was already sufficient to decide the case.

Most of the relevant statute book was built within a single decade. The Law of the Judiciary, which organizes the courts and defines the hierarchy from the first-degree courts through the appellate courts to the Supreme Court, dates from 1428H.¹⁵ The Law of Enforcement followed in 1433H,¹⁶ the Arbitration Law in the same year,¹⁷ the Law of Shari‘ah Procedure in 1435H,¹⁸ the Commercial Courts Law in 1441H,¹⁹ the Evidence Law in 1443H,²⁰ and the Civil Transactions Law in 1444H.

The order is worth noticing. The instruments governing *how* a case is heard were enacted first, over more than a decade. The instrument governing *what the parties owe each other* came last.

B. The Civil Transactions Law

The Civil Transactions Law was issued by Royal Decree No. M/191 of 29/11/1444H. It contains 721 articles covering obligations, contract, and the general principles of civil liability, and it is the first codification of that field in the Kingdom’s modern history.²¹

¹⁴Basic Law of Governance, Royal Order No. A/90 (27/8/1412H) (Saudi Arabia), arts. 46–48, 67.

¹⁵Law of the Judiciary, Royal Decree No. M/78 (19/9/1428H) (Saudi Arabia), art. 9.

¹⁶Law of Enforcement, Royal Decree No. M/53 (13/8/1433H) (Saudi Arabia).

¹⁷Arbitration Law, Royal Decree No. M/34 (18/7/1433H) (Saudi Arabia).

¹⁸Law of Shari‘ah Procedure, Royal Decree No. M/1 (22/1/1435H) (Saudi Arabia).

¹⁹CCL, *supra* note 4.

²⁰Evidence Law, Royal Decree No. M/43 (26/5/1443H) (Saudi Arabia). Its arrival repealed the entire evidence chapter of the Commercial Courts Law — articles 38 through 57 — so that evidence is now governed uniformly across the courts.

²¹CTL, *supra* note 3. The corpus used in this Article holds the full official Arabic text of all 721 articles, cross-checked article by article against the Ministry of Justice’s legal portal and, where the two diverged, against the Ministry’s

Its enactment did not displace the uncoded sources as a matter of constitutional form; the Basic Law's direction to the courts was not amended. What the code offers a Saudi judge is an available and authoritative statement, in numbered articles, of rules that previously had to be derived. Whether judges reason from it is exactly the question this Article measures, and it is a real question rather than a rhetorical one, because a judge who does not cite the code has not thereby decided unlawfully.

C. The Jurisdiction Clause

One feature of the enactment makes the measurement below more pointed than it would otherwise be.

Royal Decree No. M/191 did not only issue the Civil Transactions Law. The same decree amended article 16 of the Commercial Courts Law — the provision fixing what the commercial courts have jurisdiction to hear.²²

So the instrument that created the code, in the same breath, redrew the jurisdiction of the courts most likely to apply it. This is not a case in which a legislature enacted a code and left the courts to discover it. The legislature pointed.

The pointing also supplies this Article with a denominator that is not of its own choosing. A natural objection to any measurement of statutory uptake is that the denominator is arbitrary: the courts should not be expected to cite the whole statute book, only the part relevant to what they hear. Article 16 answers that objection in the legislature's own words rather than the author's, by defining the field the commercial courts are to occupy. Part IV reports the uptake figure under that restricted denominator as well as under the full statute book, and the restriction does not rescue the result.

own published PDF. The law is unamended: every article stands as originally enacted.

²²CCL art. 16, as amended by Royal Decree No. M/191 (29/11/1444H). Article 16 is also, by a wide margin, the most-cited single article in the entire corpus, appearing 16,265 times. The courts read their own jurisdiction clause constantly.

D. A Published Record

The Ministry of Justice publishes Saudi commercial judgments in full text through the interface its own portal uses. A published record carries the judgment; where the case was appealed, it carries the appellate decision as well, in a separate field of the same record.

Two features of the published record shape the method rather than the interpretation, and are stated here because a reader entitled to check this Article's figures needs them.

First, the interface exposes fields named for the structural parts of a judgment — its facts, its reasons, its ruling — and every one of them is null in every record. The structure of a judgment must therefore be read from its text, using the headings Saudi courts themselves write.

Second, the interface exposes a flag marking a record as an appeal. The flag is present and well-typed, and it is false on every record in the corpus, including every record that carries an appellate decision. First instance and appeal are distinguished here by the fields that hold the text, not by the flag provided for the purpose.

Neither defect affects what the courts cite. Both affect how a researcher must read the record, and a study that trusted either would report structure that is not there.

IV. Method

A reader is entitled to know how a claim about 121,207 citations was produced, and in enough detail to disbelieve it. This Part sets out the steps. It is longer than a methods note in an empirical journal would be, because the errors that matter here are legal rather than statistical: they are errors about which provision a court was talking about, and a lawyer can check them.

A. Two Corpora, Joined at the Article

The legislative side is a registry of 15,855 articles across 291 Saudi instruments, each article carrying its number, its text, its position in the instrument's chapter structure, and its status — whether it

stands as originally enacted, has been amended, or has been repealed.²³ The judicial side is 50,666 judgments published in full text.

290 of the instruments can be joined to the judgments; 106 of them are cited at least once. Of the judgments, 41,395 — 81.7 per cent — cite at least one statutory provision. The remainder cite none, which is itself worth stating: a judgment that resolves a dispute without naming a provision is not an anomaly in this record.

B. Finding a Citation

Saudi judicial prose names a provision in a recognizable shape: an ordinal or numeral for the article, a preposition, and the instrument's name — “article (16) of the Commercial Courts Law.” The shape is recognizable but not uniform, and three features of written Arabic complicate it.

Prefixes attach to the word for “article” rather than standing apart, so the same reference appears as *al-māddah*, *bil-māddah*, *lil-māddah*, *wal-māddah* and so on, each a single orthographic word.²⁴ Diacritics are optional and appear irregularly, so a pattern written against unvowelled text fails on the minority of judgments that are fully vowelled. And the article number may be a numeral in either Arabic-Indic or Western digits, or spelled as a compound ordinal.

121,207 citations are extracted. What follows is what is done with them.

C. Matching a Citation to the Provision It Names

Two steps, and the first is where the serious errors live.

The instrument. A citation names an instrument in words, and the words must be resolved to one of the 291 in the registry. String containment is the obvious method and it is wrong, because

²³Repealed articles are retained and flagged rather than deleted. A judgment given in 1441H may cite a provision repealed in 1443H, and a registry that dropped it would record that citation as unmatched, which is a different error from the one it looks like.

²⁴An early version of the pattern accepted only some of these prefixes. See *infra* Part III.F.

Saudi instrument names nest: the title of the *Implementing Regulation of the Commercial Courts Law* contains the title of the *Commercial Courts Law* entire. A matcher that accepts containment in either direction will send citations of the statute to the regulation, and it will do so silently, because every citation still matches something. Matching therefore proceeds in a fixed order: exact title first; then the longest registry title contained in the cited name; and only then the shortest registry title that contains it.²⁵

Judgments also refer anaphorically: “the Law,” “the Regulation,” “the same Law,” after naming the instrument once. 7,742 citations are of this kind, against 108,474 that name their instrument outright. Anaphora is resolved by *kind* and not by recency: “the Regulation” resolves to the most recent regulation mentioned, not to the most recent instrument of any sort. Resolving by recency sends a statute’s articles to its own implementing regulation whenever the two are discussed together, which in these judgments is constantly.

The article. The article number is then read. Compound Arabic ordinals are parsed rather than looked up, so that “the ninety-fifth after the hundred” resolves to 195. A number that resolves to an article the instrument does not contain is recorded as out of range and excluded rather than rounded to a neighbour: 1,484 citations, 1.2 per cent, fall in that class.

4.1 per cent of citations are left unmatched. They are reported as unmatched and not distributed across the instruments in proportion, because a citation whose instrument cannot be identified is not evidence about any instrument.

D. Reading Where the Citation Sits

A count of citations answers a narrower question than it appears to. A provision named in the court’s recital of what a party argued is not a provision the court applied, and treating the two alike

²⁵The consequence of getting this wrong is not marginal. In an earlier version of this analysis the containment test ran in both directions, and it misdirected roughly half of all citations — with no visible symptom, because the match rate was unaffected.

measures the bar and the bench together.

Saudi judgments carry their own internal structure and announce it with headings — the statement of the case, the reasons, the operative part — and 27,321 judgments, 53.9 per cent, can be segmented on those headings. Each citation is then located in the segment it falls in.

Segmentation alone is not enough, and the reason is a fact about Saudi judgment-writing rather than about software. The segment before the reasons is not a transcript of the parties' pleadings. It is written by the court, and a large part of it is the court narrating what it did — convening, satisfying itself of jurisdiction, ordering an expert report — with the governing provisions named as it goes. Of the 22,185 citations in that segment, at least 9,531 (43.0 per cent) sit in a sentence whose grammatical actor is the court or the circuit itself.²⁶ Treating that segment as the parties' voice would have credited litigants with the court's own procedural recitals — and, because those recitals are almost entirely procedural, would have made the bar look far more procedural than it is.

E. Choosing a Denominator

The headline figure — what share of the statute book is ever cited — depends on what counts as the statute book, and the choice can be made to produce almost any answer. It is therefore made four times, and all four are reported.

1. **The whole registry:** 15,855 articles. 11.7 per cent cited.
2. **Article 16, read broadly:** every instrument the commercial courts' jurisdiction clause reaches on a generous reading — 115 instruments, 8,356 articles. 20.6 per cent cited.
3. **Article 16, read narrowly:** the same clause read strictly — 54 instruments, 5,297 articles. 30.5 per cent cited.
4. **Instruments ever cited:** the objection conceded outright. Every instrument never cited at all is dropped from the denominator, leaving 9,049 articles in 102 instruments that the courts demonstrably do use. 20.4 per cent cited.

²⁶The attribution is a lexical cue, not an understanding of the sentence, and it is reported as a floor rather than a partition: the unattributed remainder contains both party argument and further judicial narration whose cue was not caught. On a sample of 40 attributions read by hand, 37 were correct.

The fourth is the most hostile denominator this Article can construct against itself, and the figure it yields is still under a quarter. Two of the four are drawn from the legislature’s own jurisdiction clause rather than from the author’s judgment about relevance, which is the point of Part II.C.

One robustness check belongs here. The published record contains duplicates — 796 groups of judgments sharing a text, 1,303 extra copies, 2.6 per cent of the corpus. Recomputing the headline shares over the 49,363 distinct texts alone moves the procedural share from 89.2 to 89.1 per cent. The results are not artefacts of republication.

F. Three Defects, and a Fourth Found by Accident

Each of the following produced a confident wrong answer, and none announced itself. They are reported because a reader assessing this Article’s figures should know what kind of error it is capable of, and because the fourth was found in a way worth recording.

Containment in both directions. Described above. It misdirected roughly half of all citations while leaving the match rate untouched.

Anaphora by recency. “The Law,” resolved to the last instrument named rather than the last *law* named, sent statutes’ articles to their own regulations.

The recital read as the parties’ voice. Described above. It made the central finding of Part IV.C point the wrong way.

The prefix gap. The pattern that finds a citation accepted several of the prefixes Arabic attaches to the word for “article” but not all of them, and so missed 16,682 citations outright — 15.7 per cent of everything the project had counted. It was not found by any check run against the judgments. It was found by pointing the extractor at a document written by a practising lawyer, whose prose used a prefix the corpus’s own drafting conventions had made rare.²⁷ Every figure in this Article was recomputed after the correction.

²⁷The general lesson is not about Arabic. Fifty thousand documents produced by one institution share that institution’s habits, and a tool validated only against them inherits the habits as assumptions. Input from outside the corpus tested something no amount of internal checking could reach.

Table 1: Share of the statute book cited at least once, under four denominators.

Denominator	Articles	Cited
The whole registry	15,855	11.7%
Article 16's field, read broadly	8,356	20.6%
Article 16's field, read narrowly	5,297	30.5%
Instruments the courts do cite	9,049	20.4%

V. Findings

A. Uptake at the Article Level

81.7 per cent of the judgments in this corpus cite at least one statutory provision. What they cite is a small part of what has been enacted.

Of the 15,855 articles in the registry, 1,849 — 11.7 per cent — are cited even once in 50,666 judgments. Of the 291 instruments, 106 are cited at all; 185 are never cited by any judgment in the corpus.

Table 1 reports the figure under all four denominators described in Part III.E, because the objection that the denominator is too generous is the first one a reader should make.

Two of the four denominators come from the legislature. The other two are built to be hostile to the finding: the last drops every instrument the courts never touch, so that the denominator contains only law they demonstrably use, and the share is still 20.4 per cent. Under no reading does a majority of the enacted text enter adjudication.

What is cited is also concentrated. The single most-cited article accounts for 14.6 per cent of all matched citations; the top ten account for 43.3 per cent; the top hundred for 79.2 per cent.²⁸ The applied law is not a thin, even layer over the statute book. It is a few dozen provisions, read constantly.

²⁸The most-cited article is CCL art. 16 — the jurisdiction clause — at 16,265 citations. A court reading its own jurisdiction is the corpus's most frequent act of statutory interpretation.

Table 2: Share of an instrument's own articles cited at least once.

Instrument	Articles	Cited
Commercial Courts Law	96	99.0%
Evidence Law	129	90.7%
Companies Law	281	71.2%
Bankruptcy Law	231	51.9%
Civil Transactions Law	721	27.5%

B. The Procedural Asymmetry

89.2 per cent of matched citations fall on instruments governing how a case is heard rather than what the parties owe each other.²⁹

The sharpest statement of that asymmetry compares instruments to themselves, where no denominator is chosen by anyone (Table 2).

The Commercial Courts Law has 95 of its 96 articles cited — 99.0 per cent. The Civil Transactions Law has 198 of 721 — 27.5 per cent. The first tells the court how to proceed. The second tells the court what the parties owe each other, and is the instrument the amended jurisdiction clause of the first directs the court toward.

It would be convenient to attribute this to the code's youth. That explanation is available and this Article does not exclude it; Part V returns to it. But it does not explain the ordering of the other three rows, which are not new: the Evidence Law, enacted a year before the code, is at 90.7 per cent, and the Bankruptcy Law and Companies Law, substantive instruments both, sit between. The gradient runs from procedure to substance, not from old to new.

²⁹The classification is by instrument, not by article: an instrument is procedural if its subject is the conduct of litigation — the Commercial Courts Law and its implementing regulation, the Law of Sharī'ah Procedure and its regulation, the Evidence Law, the Law of Enforcement, and the like. Recomputed over the 49,363 distinct texts alone, to remove the effect of republished judgments, the share is 89.1 per cent.

Table 3: Statutory citations by segment, across 27,321 segmented judgments.

Segment	Citations	Instruments	Procedural
Statement of the case	22,185	80	84.9%
The court's reasons	49,204	68	94.5%

C. *The Bench, Not the Bar*

The obvious objection to everything above is that the courts cite what they are given. If litigants argue procedure, judgments will be procedural, and the measurement describes the bar rather than the bench.

The objection is testable, and it fails. In the 27,321 judgments that can be segmented on their own headings — 53.9 per cent of the corpus — each citation is located in the statement of the case or in the court's reasons (Table 3).

The court's reasons are *more* procedural than the statement of the case, not less, and drawn from a narrower set of instruments: 68 against 80. In 24 instruments a statute appears in the statement of the case and never once in any court's reasons.

Attribution sharpens this rather than softening it. Of the 22,185 citations in the statement of the case, at least 9,531 — 43.0 per cent — sit in a sentence whose actor is the court itself, narrating what it did. Those are 98.6 per cent procedural, which is what one would expect of a court reciting the provisions under which it convened. The remainder, which is where the parties' arguments live, is 74.6 per cent procedural. Removing the court's own narration from the segment therefore *widens* the gap between what is put to the court and what the court answers with, from 84.9 against 94.5 per cent to 74.6 against 94.5.

Substantive instruments are raised, and are answered with procedure. Whatever is narrowing the applied law, it is not the bar.

D. Amendment Does Not Predict Use

A last explanation deserves testing, because it would make the finding uninteresting. Perhaps the uncited portion of the statute book is its dead portion — provisions superseded in substance, which the legislature has left standing and the courts have sensibly abandoned. If so, legislative attention should track judicial use: the articles a legislature has bothered to amend should be the articles courts cite.

Pooled across the 6,858 articles whose status the registry records, that is exactly what appears. The 539 amended articles are cited at 44.3 citations per article against 12.9 for the 6,319 articles standing as enacted — a ratio of 3.4 to one, and 26.0 per cent of amended articles are cited at all against 21.9 per cent of original ones.

The pooled figure is a composition effect and it does not survive a control. Amendment is not distributed evenly across instruments: the heavily amended instruments are the procedural ones, which are also the heavily cited ones, and the pooled comparison is reading that rather than any relation between amendment and use. Comparing amended against original articles *within* each instrument, across the 46 instruments where both classes appear, the amended articles are cited more often in 20, less often in 25, and equally in 1 — a sign test at $p = .551$. There is no within-instrument relation at all.

The same holds one level up. Across the 266 instruments with a recorded amendment history, of which 90 are cited, the rank correlation between how much an instrument has been amended and how much it is cited is .201 — weak, and in any case an instrument-level association that says nothing about which articles are used.

So the claim that amendment predicts citation is withdrawn, and with it the explanation it would have supported. The uncited portion of the Saudi statute book is not the portion the legislature has abandoned. It is live law that courts do not reach for.

VI. Implications and Limits

A. *For Comparative Law*

The argument of this Article is not that Saudi courts under-apply their civil code. It is that whether they do is a question with an answer, and that comparative law has been arguing a related question without ever asking this one.

Enactment and uptake are separate facts. The first is accomplished by promulgation and is visible in a gazette. The second happens, or fails to happen, in the reasoning of courts over years, and it is visible only where those courts publish. The transplant literature's disagreements — whether a borrowed rule can travel, whether it irritates its host, whether reception conditions determine later effectiveness — are all disagreements about what follows enactment.³⁰ None of them has been tested against the courts said to be applying the transplanted law, at the level of the provision, because the evidence has not existed.

It exists now, in at least one jurisdiction, and this Article's first claim is methodological: uptake is measurable, the measurement is cheap once the two corpora exist, and it is not a proxy for the thing of interest but the thing itself. A code that courts do not cite is not being applied through some invisible channel; it is not the source they are reasoning from.

The second claim is substantive and narrower, because it rests on one jurisdiction. Uptake is not merely slow — it is *shaped*. What the courts here took up first, and almost completely, is the law that tells them how to proceed. What they have taken up least is the law that tells them what the parties owe each other, which is the part a codification is normally undertaken to supply. If that pattern holds elsewhere, it suggests that 'does the transplant work?' decomposes into questions with different answers: procedural transplants may be absorbed almost immediately, because a court needs a procedure to sit at all, while substantive ones compete with whatever the court was deciding by before.

Whether it holds elsewhere is now an empirical question rather than a speculative one. The

³⁰See *supra* Part I.B.

conditions this Article needed — a machine-readable statute book, published judgments at scale, and a recent codification — are being assembled in several jurisdictions at once, including others in the Gulf. The measurement can be repeated.

B. For Legislative Evaluation

Ex post legislative evaluation has become an established practice, and it measures compliance costs, administrative burdens, and outcomes in the world. It does not, as a rule, measure whether the courts have used the text.

Uptake is an unusually cheap indicator to add. It requires no survey, no interview and no new data collection: where judgments are published and the statute book is structured, it is computed from public records, at the level of the article, and it is auditable by anyone who disagrees with it. A legislature that enacted a code three years ago can ask which of its articles have ever been applied by a court, and get a list.

The danger in such an indicator is obvious and worth stating before anyone builds a dashboard from it. **Uptake is not merit.** An article may go uncited because it is badly drafted; it may equally go uncited because it is so clear that no dispute over it reaches a court, because the conduct it governs is rare, because the disputes it governs are settled or arbitrated, or because they are heard in a court whose judgments are not in this corpus. A low uptake figure raises a question. It does not answer one, and it should never be read as a score.

Read as a diagnostic, though, it has a use no other indicator has. It identifies, by article, the parts of a statute book that adjudication has never touched, and it hands a drafting body a finite list to look at. For the Civil Transactions Law that list currently holds most of the code.

C. Limits

Five limits bound every figure above. The first two are the ones on which this Article could most plausibly be wrong.

Citation is not application. This is the deepest limit, and it is sharpened rather than softened by the jurisdiction. A Saudi judge has a source of decision that does not depend on any enacted text, and may resolve a dispute on grounds that coincide exactly with an article of the code without naming the article.³¹ Every figure here therefore measures citation, which is a *lower bound* on application. The Article's central comparison survives this — the same judges, writing in the same style, cite the procedural law almost exhaustively and the civil code barely at all, and no account of silent application explains why silence would fall so unevenly — but the absolute levels should be read as floors.

The code is four years old. The Civil Transactions Law was issued in 1444H and the corpus runs to 1448H. A code's first years are not its verdict: litigation takes time to arise under new rules, and a judgment given in 1445H may concern a contract formed before the code existed. This Article reports what uptake looks like in the years immediately following codification, which is a different question from what it will look like in 2040, and the reader should not read the second answer out of the first.

The record is what the Ministry publishes, and it is commercial. 95.0 per cent of the corpus is commercial, 85.1 per cent comes from three cities, and 87.9 per cent falls in the most recent Hijri years. That composition cuts both ways for the central finding. It is favourable in that the commercial courts are the courts whose jurisdiction clause was amended to point at the civil code, so these are the judges most directed toward it. It is unfavourable in that much civil litigation is heard elsewhere, in courts whose judgments this corpus does not contain, and the code's uptake there is simply unknown.

The extraction is lexical. 4.1 per cent of citations could not be resolved to an instrument, 1.2 per cent named an article the instrument does not contain, and 2.6 per cent of article expressions could

³¹Basic Law of Governance art. 48; *supra* Part II.A.

not be parsed at all. None is distributed in a way that would manufacture the findings, and all are reported rather than redistributed.

The segment analysis covers half the corpus. Only 53.9 per cent of judgments carry headings the segmenter can find, so the bench-and-bar result in Part V.C is measured on that half. The attribution within it is a lexical cue validated by hand on a small sample, and is reported as a floor.

VII. Conclusion

A legislature enacts a code and a gazette records it. Whether the courts then reason from it is a separate fact, and until recently an unobservable one. This Article has observed it for a jurisdiction that codified its private law within the last four years and publishes its commercial judgments in full.

The answer is that the applied law is a small, lopsided, and heavily procedural subset of the enacted law. 11.7 per cent of the statute book is ever cited, and 20.4 per cent even after every instrument the courts ignore entirely is removed from the denominator. 99.0 per cent of the articles of the law that tells the court how to proceed have been cited; 27.5 per cent of the code that tells the court what the parties owe each other. The narrowing is not the parties' doing, and it is not explained by the legislature having abandoned the unused part of the book.

Nothing here shows that the Saudi codification has failed. Four years is not a verdict, and citation is a floor on application. What it shows is that enactment and uptake are different facts, that the gap between them has a measurable shape, and that comparative law's long argument about transplanted codes has been conducted, for half a century, on the wrong side of that gap.

Data and Code

The corpus, the extraction and matching code, and the scripts that generate every figure in this Article are public.³²

³²<https://github.com/Abdullah-m7/saudi-legal-corpus-ai>. Every measurement reported here is generated into a single file of macros and typeset from it; a check refuses to build a manuscript that types one by hand.